

Marketing Analytics

Chapter 1 · Marketing Analytics as Decision-Making

Nils Wlömert

Vienna University of Economics and Business (WU)

September 2026

Before we start: who are you?

Go to **menti.com**

and enter code **XXXX XXXX**

- Your background & prior R experience
- What you want to get out of this course

QR code
(img/menti_qr.png)

Who I am...



- **Professor** of Marketing, Vienna University of Economics and Business, Institute for Retailing & Data Science, since 2021
- **Assistant Professor** of Marketing, Vienna University of Economics and Business, Institute for Interactive Marketing & Social Media, 2015–2020
- **Ph.D.** in Marketing, University of Hamburg, 2010–2014
- **Dissertation topic:** “Information Technology and Online Content Distribution”
- **Universal Music Group:** Business Analyst & Digital Marketing Manager, 2008–2010
- **Research interests:** multi-channel distribution, online content distribution, digital business models, marketing modeling, e-commerce policies, information goods, targeting & consumer privacy ...

Contact

Univ.Prof. Dr. Nils Wlömert


Room D2.1.452

nils.wloemert@wu.ac.at

T +43 1 31336 4958

sites.google.com/view/nilswloemert

Institut für Retailing & Data Science



r<<ds

Institute for Retailing
and Data Science

Follow us on LinkedIn!

[linkedin.com/company/institut-für-retailing-data-science-rds](https://www.linkedin.com/company/institut-für-retailing-data-science-rds)

Course organization

How this course works

Weekly sessions & self-study

- **Weekly in-person sessions** (see schedule) — concepts, live analysis, discussion, and lots of time for questions
- **Self-study *after* each session:** the course book on the website consolidates the session — chapter text, commented code, and a learning check; slides & PDF handouts linked from every chapter
- Each session opens with a short **paper quiz** on the previous chapter (*counts toward participation*)
- **Before session 2:** complete the self-paced **R onboarding module** (install R/RStudio *or* use Posit Cloud) — AI assistants are an endorsed study aid for setup problems

Applying the knowledge

- **Three individual assignments** transfer the methods to new business settings — reproducible reports (.Rmd), submitted via Canvas, short in-class debrief (*details in a moment*)
- **Group project:** design & conduct a real market research project — two coaching sessions, results submitted as a **pre-recorded video presentation**

Getting help

- **Forum on Canvas** first — if you have the question, so do five others; then coaching sessions & in-person sessions

Schedule — Group A (Mondays)

Course contents

Date	Time	Room	Topic
Mo 05.10.	13:00–16:00	TC.4.03	Questions, decisions & data
Mo 12.10.	13:00–17:00	D3.0.233	Describing data
Mo 19.10.	13:00–17:00	D3.0.233	Uncertainty & first comparisons
Di 27.10.	13:00–16:00	TC.5.05	Experiments I
Mo 02.11.	13:00–16:00	TC.4.03	Experiments II
Mo 09.11.	13:00–16:00	TC.4.03	Observational data I
Mo 16.11.	13:00–16:00	TC.4.03	Observational data II
Mo 23.11.	13:00–16:00	TC.4.03	Observational data III

Group project

Date	Time	Milestone
Mo 30.11.	13:00–17:00	Coaching I — questionnaire design (<i>online</i>)
December	—	Field your survey
Do 07.01.	13:00–17:00	Coaching II — data analysis (<i>online</i>)
Jan (<i>tbc</i>)	—	Video presentation: submission
Jan (<i>tbc</i>)	—	Peer review of two videos

Final exam: Di 15.12., 10:00–12:00, D5.0.001
paper & pencil, open book, in person

Assignments — Group A (Mondays)

#	Topic	Hand-out	Submission	Debrief
A1	Data handling & description in R	Mo 12.10.	Fr 23.10.	Di 27.10.
A2	Analyzing experiments	Mo 02.11.	Fr 13.11.	Mo 16.11.
A3	Regression: from correlation to causal claims	Mo 16.11.	Fr 27.11.	Mo 30.11.

- Submitted as a **reproducible report** (.Rmd → HTML) via Canvas
- Every student receives an **individual dataset** — results and conclusions differ **by design**
- **Debrief** = 3 short paper questions about *your* results at the start of the next session



Schedule — Group B (Thursdays)

Course contents

Date	Time	Room	Topic
Do 08.10.	13:00–16:00	TC.5.03	Questions, decisions & data
Do 15.10.	13:30–17:30	TC.3.03	Describing data
Do 22.10.	13:00–17:00	TC.5.15	Uncertainty & first comparisons
Do 29.10.	13:00–16:00	TC.5.03	Experiments I
Do 05.11.	13:00–16:00	TC.5.03	Experiments II
Do 12.11.	13:00–16:00	TC.5.03	Observational data I
Do 19.11.	13:00–16:00	TC.5.03	Observational data II
Do 26.11.	13:30–16:30	TC.3.03	Observational data III

Group project

Date	Time	Milestone
Do 03.12.	14:00–18:00	Coaching I — questionnaire design (<i>online</i>)
December	—	Field your survey
Fr 08.01.	14:00–18:00	Coaching II — data analysis (<i>online</i>)
Jan (<i>tbc</i>)	—	Video presentation: submission
Jan (<i>tbc</i>)	—	Peer review of two videos

Final exam: Di 15.12., 10:00–12:00, D5.0.001
paper & pencil, open book, in person

Assignments — Group B (Thursdays)

#	Topic	Hand-out	Submission	Debrief
A1	Data handling & description in R	Do 15.10.	Mo 26.10.	Do 29.10.
A2	Analyzing experiments	Do 05.11.	Mo 16.11.	Do 19.11.
A3	Regression: from correlation to causal claims	Do 19.11.	Mo 30.11.	Do 03.12.

- Submitted as a **reproducible report** (.Rmd → HTML) via Canvas
- Every student receives an **individual dataset** — results and conclusions differ **by design**
- **Debrief** = 3 short paper questions about *your* results at the start of the next session



Our AI policy

AI assistants are welcome in the assignments — they may support your work at every step, but not do it for you: a submission without individual contribution earns no credit, disclosed or not. You must be able to **explain everything you submit** (that is what the debrief checks), and each report ends with a short **AI-use disclosure**. Honest disclosure is never penalized; concealing substantial AI use is. The final exam is the unassisted layer: paper & pencil, in person.

***“How much AI is allowed?” — As much as you like, provided you can stand here and explain what was done and why.**

Grading

Component	Weight
Market research group project (design → data collection → analysis → video presentation)	30%
Individual take-home assignments (3 × reproducible report + in-class debrief)	20%
Final exam (concepts & methods — no coding · course slides may be used , paper & pencil, in person)	40%
Class participation (weekly quizzes · debriefs · video peer review · in class & forum)	10%

- Pass threshold: **60%** weighted total
- Group project: **peer assessment + contribution & AI-use statement** enter the individual grade; **every member presents on camera**
- Participation = **more than showing up**: quizzes, debriefs & peer review are the documented base — the rest comes from **active contributions** (quality over quantity)
- Assignment dates: see slides · all details on the course site

Grading scheme

Overall percentage	Grade
90–100%	1 (excellent)
80–89%	2 (good)
70–79%	3 (satisfactory)
60–69%	4 (sufficient)
< 60%	5 (fail)

To pass, your **weighted final grade** must reach at least **60%** — across all components combined; there is no minimum per component.

Course materials


- **Main reference: the course website** — full chapters with explanations, commented R code, and learning checks; **slides & PDF handouts** linked from every chapter
- Everything is **free & always current** — no textbook purchase required
- **Optional further reading** (linked on the site): *Data Analysis for Business, Economics, and Policy* (Békés & Kézdi) · *R for Data Science* (Wickham) · and a curated list per topic
- **Canvas**: assignment hand-outs & submissions, forum, announcements
- Example exams: provided on the site **before the exam**

Contacts



Lecture & group project

Univ.Prof. Dr. Nils Wlömert
nils.wloemert@wu.ac.at



Assignments

XXX
XXX@wu.ac.at

Before you write an email: course site → forum on Canvas → ask in class. If you have the question, so do five others.

The group project in one sentence

You will **design a survey containing a real randomized experiment**, field it, analyze the data, and present the results — everything this course teaches, built by you.

(Details follow — and at the end of today's session you will see why “randomized” is the most important word in that sentence.)

The group project: two parts

Part 1 · Questionnaire design & data collection

- Pick a topic: **14 proposed topics** (with starter literature) — or suggest your own
- Design a questionnaire (≈ **30 questions**, ~5 per member) in **Qualtrics** — free account provided
- Your design contains a **real experiment**: participants randomly assigned to conditions
- **Coaching I**: present your draft (research problem, design, measurement, planned analyses) → feedback → revise
- **Field it**: each member recruits ≈ **20 respondents** → ≈ **120 responses** per group

Part 2 · Data analysis & presentation

- Import & clean your Qualtrics data in R
- Analyze: descriptives → your **experimental comparison** → heterogeneity (for whom does it work?)
- **Coaching II**: work through your analysis with us (*online*)
- Submit a **.zip**: pre-recorded **video presentation** + data + reproducible R code + slides
- Afterwards: **peer review** — you watch & comment on two other groups' videos

Individual grades: peer assessment within the group + a short **contribution & AI-use statement** + every member presents a substantive part on camera.

Group project: what makes it good

- A **precise research question** with managerial relevance — and hypotheses stated *before* data collection
- A clean **experimental manipulation**: everything constant except the treatment; randomizer placed correctly
- **Deliberate measurement**: a dependent variable with real variation (e.g., willingness to pay), likely confounders measured
- Analyses **matched to the design** — justified test choice, honest uncertainty, assumptions checked
- A presentation that **tells the decision-maker what to do** — not a parade of charts

Full guidelines, topic list, Qualtrics templates & how-tos: [group project chapter](#) on the course site.

Course goals

After taking this course, you should...

- ...know why analytical skills have become **career-defining** for marketers — and which skills stay scarce when AI does the mechanics
- ...know **what data can and cannot tell you** — and recognize which question you are actually asking: *describe, predict, explain, decide*
- ...know **when to use which method** — and which assumptions the answer rests on
- ...be able to **implement the methods in R** — with AI assistants as accelerators, not replacements
- ...be able to **judge analyses made by others** — a colleague's slide deck, a platform dashboard, an AI-generated report
- ...be able to **conduct your own research project** — from question to data to recommendation
- ...no longer be **afraid of numbers & statistics** (in case you were)
- ...feel the **desire to learn more** about analytics

Today's session

Part A — Why we are here

- 1.1 Why marketing runs on data
- 1.2 A \$50 million question

Part B — The framework of the course

- 1.3 Four kinds of questions
- 1.4 Correlation, causation & what data can show
- 1.5 The cast of variables
- 1.6 Why randomization works
- 1.7 Between the poles: quasi-experiments

Part C — Making it work


- 1.8 From business problem to research design
- 1.9 The analytics workflow (and AI assistants)
- 1.10 The course at a glance

One promise for today: no formulas, no R output, no tests. Concepts and pictures only — the machinery follows in later sessions.

1.1 Why marketing runs on data

Marketing is reorganizing around AI and data

- AI already performs **about a quarter** of marketing activities — nearly **doubled** in two years
- Generative AI alone: from 7.0% to 22.4% of activities (+220% in two years)
- Marketing leaders expect **more than half** of marketing work to be performed by AI **within three years**



... while budgets are flat: every euro must justify itself

1.7%

growth of total marketing spending
— the **weakest since 2021**

15.3%

of marketing budgets now go to AI
initiatives

78.2%


name “**higher ROI per marketing dollar**” as the #1 reason to build marketing capabilities

“CMOs are being asked to deliver growth, efficiency and transformation **without meaningful budget expansion**. Those who succeed will make deliberate, **data-driven trade-offs**.”

— Gartner CMO Spend Survey (2026)


Sources: The CMO Survey (2026); Gartner CMO Spend Survey (2026), n = 401 CMOs, North America & Europe.

Who turns analytics into decisions?



Source: McKinsey — “Analytics translator: the new must-have role” (Henke, Levine & McInerney, McKinsey Quarterly).

Who turns analytics into decisions?



Source: McKinsey — “Analytics translator: the new must-have role” (Henke, Levine & McLnerney, McKinsey Quarterly).


The scarce skill is judgment

The ambition–readiness gap (Gartner 2026):

- 70% of CMOs: becoming an AI leader is a critical goal
- ... but only 30% are ready to **scale** AI capabilities

The labor market (WEF Future of Jobs 2025):

- **Analytical thinking** = the #1 core skill — 7 of 10 employers consider it essential
- **AI & big data** = the fastest-growing skill for 2025–2030




Source: World Economic Forum, Future of Jobs Report 2025

The mechanical layer of marketing is being automated. What remains scarce — and what this course trains — is **judgment**: knowing which analysis answers which managerial question, and whether the resulting number can be trusted.

1.2 A \$50 million question

eBay, 2012: does paid search advertising pay off?

- eBay spent ~ \$50M **per year** on paid search ads (Google, Bing, Yahoo)
- The analytics team examined the data: users **exposed to ads bought far more** than users who weren't
- Measured “returns” on brand-keyword ads looked spectacular



How the “returns” were measured: click attribution

1. A user types “**ebay guitar**” into Google. The top result is eBay’s **paid** link; directly beneath it sits eBay’s **organic** link.
2. She clicks the paid link. The URL carries a **tracking parameter**, so eBay’s systems log: this visit arrived via a paid search click, on this keyword, at this cost-per-click.
3. Within the **attribution window** (say, some days), she buys a €150 guitar. The attribution system **credits the €150 to paid search** — the sale is “attributed” because it followed an ad click.
4. Aggregate over millions of users: **attributed revenue ÷ ad spend** = a spectacular return, per keyword, in a dashboard, updated daily.

Discuss (5 min): Would you base the decision to invest in keyword advertising on this analysis? What could be potential problems with it?

What the experiment revealed

- **Brand keywords** (“ebay shoes”): ads halted entirely on Yahoo! & Microsoft, Google kept as comparison → **~99.5% of the paid clicks simply returned via the free organic link below the ad**
- **Non-brand keywords** (“used guitar”): all ads suspended in a **randomly selected ~30% of US regions for 60 days**, change in sales compared against the other regions
- Effect only among **new & infrequent users**; frequent shoppers — most of the clicks and cost — were unmoved
- The naive analysis had credited the shoppers’ **intent** to the ads; true average returns: **negative**

Smart people, huge data, careful computation — and a conclusion that was off by tens of millions of dollars. **The problem was not the math. It was the comparison.**

The comparison that matters

The causal question: *What would eBay's sales have been **without** the ads — everything else unchanged?*

	Observed	Needed
Group A	Exposed users' sales	Exposed users' sales
Group B	Unexposed users' sales	The same users , unexposed
Difference measures...	Exposure + intent	Effect of exposure only

The “needed” Group B is a **counterfactual** — a world we can never observe directly. All of causal inference is about **constructing a credible stand-in** for it. This single idea

organizes the entire course.

1.3 Four kinds of questions

Every analytics task answers one of four questions

#	Question	Type	Marketing examples
1	What is going on?	Descriptive	Segment sizes, churn rate, funnel drop-off, satisfaction tracking
2	What will happen?	Predictive	Which customers will churn next quarter? Demand next month?
3	What happens if we act?	Causal	Does the discount <i>cause</i> extra sales? Did the campaign work?
4	What should we do?	Prescriptive	Set the price; allocate budget across channels

Prescription requires causation. To choose an action you must know what each action would *cause* — i.e., compare counterfactuals. Prediction alone is not enough... and confusing 2 with 3 is precisely the eBay mistake.

Prediction $\neq$ causation: same data, different questions

“Customers who see our newsletter churn 40% less.”

As a predictive statement ✓

Newsletter readership *predicts* retention — useful for **targeting**: flag non-readers as churn risks.

Valid even if the newsletter does nothing.

As a causal statement ✗ (so far)


“Send more newsletters $\rightarrow$ churn drops”?
Only if readership *causes* retention — but engaged customers **both** read newsletters **and** stay.

Requires a counterfactual comparison.

Classify these (2 min, neighbor discussion): (a) “Which ad creative will get the highest CTR?” (b) “How large is our Gen-Z segment?” (c) “Should we raise the price by 5%?”

1.4 Correlation, causation & what data can show

Two pictures of “X relates to Y”



The two pictures can look **equally convincing** — yet only (a) licenses a causal conclusion. The difference is invisible in the plot: it lies in **how the data came to be** (the *data-generating process*).

Three ways to get the same scatter plot

Reverse causality: who is causing whom?

Observed: regions with higher ad spend show higher sales.



Story 1: advertising drives sales.

Story 2: budgets are **set as a % of sales** — sales drive ad spend.

Both stories produce identical data. Marketing budget rules make reverse causality a **routine** hazard in marketing data, not an exotic one.

1.5 The cast of variables


The two lead roles



- **Outcome / dependent variable (DV):** the quantity the manager cares about
- **Treatment / independent variable (IV):** the lever the manager can pull



Defining these two precisely — *whose* purchases, over *what* horizon, *which* discount — is half of research design (§1.8).

The villain: the confounder



- Causes **both** treatment and outcome
- Creates correlation **without** causation — panel 2 of our gallery
- **The eBay confounder:** purchase intent → people search for “eBay” → see the ad → and buy anyway
- Whole toolkits (Ch. 5–6) exist to deal with confounders in observational data

Supporting roles: mediator and moderator



Mediator — the *how*


Lies **on the path**: the mechanism through which X works.
Why care? It tells you **what to optimize** (creative that builds awareness).

Moderator — the *for whom / when*

Points **at the arrow**: changes the **strength** of the effect.
“The discount works for new customers, not loyal ones.”


Today: **recognize** the roles. Chapter 5 teaches how to **handle** each one — and shows that mishandling them (e.g., “controlling” for a mediator) can wreck an analysis.

The full cast around one causal question



1.6 Why randomization works

Randomization cuts every arrow into the treatment




- Assignment now depends **only on the coin** — not on intent, engagement, income, mood, ...
- Confounders (known **and unknown**) are balanced across groups *in expectation*
- The control group becomes a valid **counterfactual**: same kinds of people, minus the treatment
- Differences in outcomes can then be attributed to the treatment (up to chance — quantified in Ch. 3)

“Balanced — even on what we never measured”

Covariate balance: exposed vs. unexposed group (simulated, n = 2,000)

Standardized mean difference per variable · dashed lines: conventional ± 0.1 band



Left: the groups differ systematically — including on **intent**, which we (realistically) never observe. Right: one coin flip balances everything at once, **without knowing what needs balancing**. That is the magic randomization buys.

The price of observational data

Causal conclusions from observational data require assuming:

No unobserved confounders — every variable that drives both treatment and outcome has been measured and correctly adjusted for.

- This assumption is **untestable in principle**: you cannot check balance on variables you did not observe
- Experiments **buy** the assumption *by design*; observational studies must **argue** for it — with domain knowledge, not statistics
- This asymmetry is the most important sentence of the course: it determines what your data can and cannot claim
- It is *not* a counsel of despair — Ch. 5–6 develop the toolkit for making the argument as credible as possible

1.7 Between the poles: quasi-experiments

Sometimes the world randomizes for you (almost)

A spectrum of designs:

	Experiment	Quasi-experiment	Observational
Who assigns treatment?	The researcher (coin flip)	External event / rule	The units themselves
Example	A/B test of a discount	Platform launches in some countries first	Newsletter readers vs. non-readers
Key requirement	Correct execution	“As-good-as-random” timing/rule	No unobserved confounders

- Policy changes, staggered rollouts, regional boundaries → treatment variation **not chosen by the units**
- Workhorse method: **difference-in-differences** — compare *changes* in treated vs. untreated markets
- Extremely common in practice (and in marketing research) → dedicated session in **Module 6 (panel data)**

1.8 From business problem to research design

The translation template

Manager: “*I want to know whether our loyalty program is worth the money.*”

Design element	Your translation
Outcome (DV)	Customer contribution margin, 12 months
Treatment (IV)	Loyalty program membership
Unit of analysis	Customer
Question type	Causal (→ prescriptive decision)
Likely confounders	Purchase frequency & engagement <i>before</i> joining
Feasible design?	Randomize invitations? Staggered regional rollout? Observational + adjustment?

A precise, falsifiable formulation (“membership increases 12-month margin”) beats a vague aspiration (“the program creates value”) — it dictates data, design, and analysis.

Your turn (pairs, 10 min)

Translate one of these into the template:

1. “Should we invest more in influencer marketing?”
2. “Is our app driving retention, or do loyal customers just download the app?”
3. “Did last spring’s rebranding damage sales in Germany?”

For each: DV, IV, unit, question type, two plausible confounders, and the **best feasible design** (experiment / quasi-experiment / observational).

1.9 The analytics workflow — and AI assistants

Where AI helps — and where it cannot replace you

Workflow step	Nature of the work	AI assistant...
Formulate the question	Judgment	can challenge you — cannot know your business
Choose design & method	Judgment	suggests options — <i>you</i> own the assumptions
Write & debug R code	Mechanical	excellent — use it freely, including in this course
Interpret results	Judgment	drafts text — but cannot verify the causal claim
Communicate to stakeholders	Mixed	strong support, needs your framing

Course policy in one line: use AI to *accelerate* your learning of the judgment steps, not to *skip* them. Assessments in this course are designed so that the judgment steps are what earns the points.

1.10 The course at a glance

One map for the whole course

Module	Sessions	Driving question
1 · Questions, decisions & data	05.10.	Which question am I really asking? (<i>today</i>)
2 · Describing data	12.10.	What is going on?
3 · Uncertainty	19.10.	How sure can I be?
4 · Learning from experiments	19.10. / 27.10.	Did it work?
5 · Learning from observational data	02.11. / 09.11. / 16.11.	What if we cannot randomize?
6 · Panel data & quasi-experiments	30.11.	What if the world randomizes for us?
7 · From estimate to decision	07.01.	What should we do?
Group project coaching	23.11. / part of 07.01.	Your own experiment, end to end
Final exam (in person)	15.12.	—

Group project preview: design a survey **containing a real randomized experiment**, field it, analyze it. Everything from today — counterfactuals, randomization, confounders — becomes something you *build*, not just recognize.

Learning check — can you spot the reasoning?

1. An analyst reports: “*App users have 2× the lifetime value of non-users — the app clearly pays off.*” Which question type is being **answered**, which is being **claimed**? Name the most plausible confounder.
2. A platform’s advertising dashboard compares purchases of users who clicked an ad vs. users who saw it and didn’t click. Is this an experiment? What role does “interest in the product” play?
3. Your CMO wants to know whether to extend a price promotion nationally after a pilot in Vienna. What design element makes the pilot informative — and what is the counterfactual?
4. For each: newsletter → (*opens* →) purchases; discount effect stronger for students — mediator or moderator?

Next session

Module 2 · Describing data — *What is going on?*

- Getting data into R, cleaning it, and not trusting it too early
- Summary statistics and the visualizations that earn their place
- **Before the session:** complete the R onboarding module (installation *or* Posit Cloud) — link on the course site. Any AI assistant is an endorsed study aid for this.

Course book chapter 1 (consolidation reading for today) is online — including all simulated examples from these slides, with their code.