

# Marketing Analytics

Chapter 1 · Marketing Analytics as Decision-Making

Nils Wlömert

Vienna University of Economics and Business (WU)

September 2026

# Who I am...



- **Professor** of Marketing, Vienna University of Economics and Business, Institute for Retailing & Data Science, since 2021
- **Assistant Professor** of Marketing, Vienna University of Economics and Business, Institute for Interactive Marketing & Social Media, 2015–2020
- **Ph.D.** in Marketing, University of Hamburg, 2010–2014
- **Dissertation topic:** “Information Technology and Online Content Distribution”
- **Universal Music Group:** Business Analyst & Digital Marketing Manager, 2008–2010
- **Research interests:** multi-channel distribution, online content distribution, digital business models, marketing modeling, e-commerce policies, information goods, targeting & consumer privacy ...

## Contact

Univ.Prof. Dr. Nils Wlömert

Room D2.1.452

[nils.wloemert@wu.ac.at](mailto:nils.wloemert@wu.ac.at)

T +43 1 31336 4958

[sites.google.com/view/nilswloemert](https://sites.google.com/view/nilswloemert)

# Institut für Retailing & Data Science



r<<ds

Institute for Retailing  
and Data Science

Follow us on LinkedIn!

[linkedin.com/company/institut-für-retailing-data-science-rds](https://www.linkedin.com/company/institut-für-retailing-data-science-rds)

# Before we start: who are you?

Go to **menti.com**

and enter code **XXXX XXXX**

- Your background & prior R experience
- What you want to get out of this course

*QR code*  
*(img/menti\_qr.png)*

# Course organization

# How this course works

## Weekly sessions & self-study

- **Weekly in-person sessions** (see schedule) — concepts, live analysis, discussion, and lots of time for questions
- **Self-study *after* each session:** the course book on the website consolidates the session — chapter text, commented code, and a learning check; slides, code & PDF handouts linked from every chapter
- Each session opens with a short **paper quiz** on the previous chapter (*counts toward participation*)
- **Before session 2:** complete the self-paced **R onboarding module** (install R/RStudio *or* use Posit Cloud) — AI assistants are an endorsed study aid for setup problems

## Applying the knowledge

- **Three individual assignments** transfer the methods to new business settings — reproducible reports (.Rmd), submitted via Canvas, short in-class debrief (*details in a moment*)
- **Group project:** design & conduct a real market research project — two coaching sessions, results submitted as a **pre-recorded video presentation**

## Getting help

- **Forum on Canvas / AI-Tools** first; then coaching sessions & in-person sessions

# Schedule — Group A (Mondays)

## Course contents

| Date      | Time        | Room     | Topic                           |
|-----------|-------------|----------|---------------------------------|
| Mo 05.10. | 13:00–16:00 | TC.4.03  | Questions, decisions & data     |
| Mo 12.10. | 13:00–17:00 | D3.0.233 | Describing data                 |
| Mo 19.10. | 13:00–17:00 | D3.0.233 | Uncertainty & first comparisons |
| Di 27.10. | 13:00–16:00 | TC.5.05  | Experiments I                   |
| Mo 02.11. | 13:00–16:00 | TC.4.03  | Experiments II                  |
| Mo 09.11. | 13:00–16:00 | TC.4.03  | Observational data I            |
| Mo 16.11. | 13:00–16:00 | TC.4.03  | Observational data II           |
| Mo 23.11. | 13:00–16:00 | TC.4.03  | Observational data III          |

## Group project

| Date               | Time        | Milestone                                           |
|--------------------|-------------|-----------------------------------------------------|
| Mo 30.11.          | 13:00–17:00 | Coaching I — questionnaire design ( <i>online</i> ) |
| December           | —           | Field your survey                                   |
| Do 07.01.          | 13:00–17:00 | Coaching II — data analysis ( <i>online</i> )       |
| Jan ( <i>tbc</i> ) | —           | Video presentation: submission                      |
| Jan ( <i>tbc</i> ) | —           | Peer review of two videos                           |

**Final exam: Di 15.12., 10:00–12:00, D5.0.001**  
*paper & pencil, open book, in person*

# Assignments — Group A (Mondays)

| #  | Topic                                         | Hand-out  | Submission | Debrief   |
|----|-----------------------------------------------|-----------|------------|-----------|
| A1 | Data handling & description in R              | Mo 12.10. | Fr 23.10.  | Di 27.10. |
| A2 | Analyzing experiments                         | Mo 02.11. | Fr 13.11.  | Mo 16.11. |
| A3 | Regression: from correlation to causal claims | Mo 16.11. | Fr 27.11.  | Mo 30.11. |

- Submitted as a **reproducible report** (.Rmd → HTML) via Canvas
- Every student receives an **individual dataset** — results and conclusions differ **by design**
- **Debrief** = 3 short paper questions about *your* results at the start of the next session



# Schedule — Group B (Thursdays)

## Course contents

| Date      | Time        | Room    | Topic                           |
|-----------|-------------|---------|---------------------------------|
| Do 08.10. | 13:00–16:00 | TC.5.03 | Questions, decisions & data     |
| Do 15.10. | 13:30–17:30 | TC.3.03 | Describing data                 |
| Do 22.10. | 13:00–17:00 | TC.5.15 | Uncertainty & first comparisons |
| Do 29.10. | 13:00–16:00 | TC.5.03 | Experiments I                   |
| Do 05.11. | 13:00–16:00 | TC.5.03 | Experiments II                  |
| Do 12.11. | 13:00–16:00 | TC.5.03 | Observational data I            |
| Do 19.11. | 13:00–16:00 | TC.5.03 | Observational data II           |
| Do 26.11. | 13:30–16:30 | TC.3.03 | Observational data III          |

## Group project

| Date               | Time        | Milestone                                           |
|--------------------|-------------|-----------------------------------------------------|
| Do 03.12.          | 14:00–18:00 | Coaching I — questionnaire design ( <i>online</i> ) |
| December           | —           | Field your survey                                   |
| Fr 08.01.          | 14:00–18:00 | Coaching II — data analysis ( <i>online</i> )       |
| Jan ( <i>tbc</i> ) | —           | Video presentation: submission                      |
| Jan ( <i>tbc</i> ) | —           | Peer review of two videos                           |

**Final exam: Di 15.12., 10:00–12:00, D5.0.001**  
*paper & pencil, open book, in person*

# Assignments — Group B (Thursdays)

| #  | Topic                                         | Hand-out  | Submission | Debrief   |
|----|-----------------------------------------------|-----------|------------|-----------|
| A1 | Data handling & description in R              | Do 15.10. | Mo 26.10.  | Do 29.10. |
| A2 | Analyzing experiments                         | Do 05.11. | Mo 16.11.  | Do 19.11. |
| A3 | Regression: from correlation to causal claims | Do 19.11. | Mo 30.11.  | Do 03.12. |

- Submitted as a **reproducible report** (.Rmd → HTML) via Canvas
- Every student receives an **individual dataset** — results and conclusions differ **by design**
- **Debrief** = 3 short paper questions about *your* results at the start of the next session




# Our AI policy

**AI assistants are welcome in the assignments** — they may support your work at every step, but not do it for you: a submission without individual contribution earns no credit, disclosed or not. You must be able to **explain everything you submit** (that is what the debrief checks), and each report ends with a short **AI-use disclosure**. Honest disclosure is never penalized; concealing substantial AI use is. The final exam is the unassisted layer: paper & pencil, in person.

**\*“How much AI is allowed?” — As much as you like, provided you can stand here and explain what was done and why.**

# Use it — but don't let it study for you



- **26,811 Chinese students, 30-month panel: AI adoption → homework +18%, time -30% — but exam scores -20%; entrance exams -18/-24% after two years**
- **Driven by the ~80% who outsource; students who used AI but spent the same time on homework: no learning loss**

Sources: Strömberg, Lei & Wu (2026); Autor et al. (2026, NBER WP 35720); Bastani et al. (2025, PNAS); Kestin et al. (2025). Figures: Strömberg et al. (2026), Fig. 2 — red x: AI adopters, blue o: never-AI students.

# Grading

| Component                                                                                                         | Weight |
|-------------------------------------------------------------------------------------------------------------------|--------|
| Market research <b>group project</b> (design → data collection → analysis → <b>video presentation</b> )           | 30%    |
| Individual <b>take-home assignments</b> (3 × reproducible report + <b>in-class debrief</b> )                      | 30%    |
| <b>Final exam</b> (concepts & methods — no coding · <b>course slides may be used</b> , paper & pencil, in person) | 30%    |
| <b>Class participation</b> (weekly quizzes · debriefs · video peer review · in class & forum)                     | 10%    |

- Pass threshold I: **60%** weighted total
- Pass threshold II: you need to be present **80%** of the in-class time
- Group project: **peer assessment + contribution & AI-use statement** enter the individual grade
- Participation = **more than showing up**: quizzes, debriefs & peer review are the documented base — the rest comes from **active contributions** (quality over quantity)
- Assignment dates: see slides · all details on the course site

# Grading scheme

| Overall percentage | Grade            |
|--------------------|------------------|
| 90–100%            | 1 (excellent)    |
| 80–89%             | 2 (good)         |
| 70–79%             | 3 (satisfactory) |
| 60–69%             | 4 (sufficient)   |
| < 60%              | 5 (fail)         |

To pass, your **weighted final grade** must reach at least **60%** — across all components combined; there is no minimum per component.

# Course materials

- **Main reference: the course website** — full chapters with explanations, commented R code, and learning checks; **slides & PDF handouts** linked from every chapter
- Everything is **free & always current** — no textbook purchase required
- **Optional further reading** linked on the site
- **Canvas:** assignment hand-outs & submissions, forum, announcements
- Example exams: provided on Canvas **before the exam**

# Contacts



## Lecture & group project

Univ.Prof. Dr. Nils Wlömert  
[nils.wloemert@wu.ac.at](mailto:nils.wloemert@wu.ac.at)



## Assignments

Anna Di Donato  
[anna.di.donato@wu.ac.at](mailto:anna.di.donato@wu.ac.at)

**Before you write an email:** course site → AI tools → forum on Canvas → ask in class. If you have the question, so do five others.

# The group project in one sentence

You will **design a survey containing a real randomized experiment**, field it, analyze the data, and present the results — everything this course teaches, built by you.

*(Details follow — and at the end of today's session you will see why “randomized” is the most important word in that sentence.)*

# The group project: two parts

## Part 1 · Questionnaire design & data collection

- Pick a topic: **14 proposed topics** (with starter literature) — or suggest your own
- Design a questionnaire (≈ **30 questions**, ~5 per member) in **Qualtrics** — free account provided
- Your design contains a **real experiment**: participants randomly assigned to conditions
- **Coaching I**: present your draft (research problem, design, measurement, planned analyses) → feedback → revise
- **Field it**: each member recruits ≈ **20 respondents** → ≈ **120 responses** per group

## Part 2 · Data analysis & presentation

- Import & clean your Qualtrics data in R
- Analyze: descriptives → your **experimental comparison** → heterogeneity (for whom does it work?)
- **Coaching II**: work through your analysis with us (*online*)
- Submit a **.zip**: pre-recorded **video presentation** + data + reproducible R code + slides
- Afterwards: **peer review** — you watch & comment on two other groups' videos

**Individual grades:** peer assessment within the group + a short **contribution & AI-use statement**.

# Group project: what makes it good

- A **precise research question** with managerial relevance — and hypotheses stated *before* data collection
- A clean **experimental manipulation**: everything constant except the treatment; randomizer placed correctly
- **Keep it simple**: 2 conditions are sufficient (your sample size will be limited)
- **Deliberate measurement**: a dependent variable with real variation (e.g., willingness to pay)
- Incorporate **moderators** to analyse heterogeneous treatment responses
- Incorporate **mediators** to understand the process
- Analyses **matched to the design** — justified test choice, honest uncertainty, assumptions checked
- A presentation that **tells the decision-maker what to do** — not a parade of charts

Full guidelines, topic list, Qualtrics templates & how-tos: [group project chapter on the course site](#).

# Course goals

After taking this course, you should...

- ...know why analytical skills have become **career-defining** for marketers — and which skills stay scarce when AI does the mechanics
- ...know **what data can and cannot tell you** — and recognize which question you are actually asking: *describe, predict, explain, decide*
- ...know **when to use which method** — and which assumptions the answer rests on
- ...be able to **implement the methods in R** — with AI assistants as accelerators, not replacements
- ...be able to **judge analyses made by others** — a colleague's slide deck, a platform dashboard, an AI-generated report
- ...be able to **conduct your own research project** — from question to data to recommendation
- ...no longer be **afraid of numbers & statistics** (in case you were)
- ...feel the **desire to learn more** about analytics

# Today's session

## Part A — Why we are here

- 1.1 Why marketing runs on data
- 1.2 A \$50 million question

## Part B — The framework of the course

- 1.3 Four kinds of questions
- 1.4 Correlation, causation & what data can show
- 1.5 The cast of variables
- 1.6 Why randomization works
- 1.7 Between the poles: quasi-experiments

## Part C — Making it work


- 1.8 From business problem to research design
- 1.9 The analytics workflow (and AI assistants)
- 1.10 The course at a glance

**One promise for today:** no formulas, no R output, no tests. Concepts and pictures only — the machinery follows in later sessions.

# 1.1 Motivation: data and AI in marketing

# Marketing is reorganizing around AI and data

- AI already performs **about a quarter** of marketing activities — nearly **doubled** in two years
- Generative AI alone: from 7.0% to 22.4% of activities (+220% in two years)
- Marketing leaders expect **more than half** of marketing work to be performed by AI **within three years**



# ... while budgets are flat: every euro must justify itself

**1.7%**

growth of total marketing spending  
— the **weakest since 2021**

**1.7%**

of marketing budgets now go to **AI**  
**initiatives**

**1.7%**

name “**higher ROI per marketing dollar**” as the #1 reason to build marketing capabilities

“CMOs are being asked to deliver growth, efficiency and transformation **without meaningful budget expansion**. Those who succeed will make deliberate, **data-driven trade-offs**.”

— Gartner CMO Spend Survey (2026)

Sources: The CMO Survey (2026); Gartner CMO Spend Survey (2026), n = 401 CMOs, North America & Europe.


# The scarce skill is judgment

## The ambition–readiness gap (Gartner 2026):

- 70% of CMOs: becoming an AI leader is a critical goal
- ... but only 30% are ready to scale AI capabilities

## The labor market (WEF Future of Jobs 2025):


- Analytical thinking = the #1 core skill — 7 of 10 employers consider it essential
- AI & big data = the fastest-growing skill for 2025–2030



Source: World Economic Forum, Future of Jobs Report 2025


The mechanical layer of marketing is being automated. What remains scarce is judgment: knowing which analysis answers which managerial question, and whether the resulting number can be trusted.

# Who turns analytics into decisions?



Source: McKinsey — “Analytics translator: the new must-have role” (Henke, Levine & McInerney, McKinsey Quarterly).

# Who turns analytics into decisions?




Source: McKinsey — “Analytics translator: the new must-have role” (Henke, Levine & McLnerney, McKinsey Quarterly).

# 1.2 Case study: search advertising at eBay

# eBay, 2012: Does paid search advertising pay off?

Figure 1: Google Ad Examples



Blake, Nosko & Tadelis (2015), Fig. 1 — Google results pages for the searches “used gibson les paul” and “macys”.


# eBay, 2012: what the dashboards showed

eBay spent ~ \$50M per year on paid search ads (Google, Bing, Yahoo). Here is how the return on this spending was measured:

1. A user types “**ebay guitar**” into Google and clicks eBay’s **paid ad** at the top of the results page
2. The **tracking parameter** in the URL logs her visit as *arrived via paid search* (with keyword & cost-per-click)
3. All other visitors are logged with a different **referrer** (organic search, other sites, direct) — and serve as the **comparison group**
4. Purchases within the **attribution window** are credited to the arrival channel → the picture on the right, and with it: **attributed revenue ÷ ad spend**, per keyword, updated daily


## What the analysts saw

Purchases per user, by channel of arrival (simulated)



**Discuss (5 min):** As eBay’s CMO, would you sign off on next year’s \$50M based on this evidence? If yes — what convinces you? If not — **how would you measure the payoff instead?**

# What happened when the ads were switched off



Paid clicks fell to zero — and **~99.5% of that traffic instantly reappeared via the free organic link**. Total visits to eBay: essentially unchanged.

Blake, Nosko & Tadelis (2015), Fig. 2 — brand-keyword ads halted on MSN (March 2012) and Google (July–August 2012); solid line: paid clicks, dotted: organic clicks.

# Reading the fine print

Brand keywords (“ebay”, “ebay shoes”) — the platform halts:

- The users clicking the ad were **already on their way to eBay** — the ad intercepted traffic the free organic link would have caught (~99.5% substitution)
- Attributed revenue: almost entirely **re-labeled, not created** → true effect of brand ads ≈ **zero**


Non-brand keywords (“used guitar”) — the regional experiment:

- Ads suspended in a **randomly selected ~30% of US regions for 60 days**; **total sales** compared against the remaining regions — a *real* comparison group, created by design
- Sales moved **only among new and infrequent users; frequent shoppers** — most of the clicks and most of the ad cost — were unmoved
- **The customers most likely to click were the customers least influenced** → averaged over everyone, returns on the spending: **negative**


Smart people, huge data, careful computation — and a conclusion off by tens of millions of dollars. **The problem was not the math. It was the comparison.**

Blake, Nosko & Tadelis (2015), *Econometrica* — details in the course book.

# Why the dashboard and the experiment disagree



# Why the dashboard and the experiment disagree



**Purchase intent drives both** who sees the ads *and* who buys. The dashboard's correlation mixes the ad effect with the intent difference — and in this case, intent was nearly the whole story.

# The comparison that matters

The causal question: *What would eBay's sales have been **without** the ads — everything else unchanged?*

|                            | What eBay observed                       | What the decision needs            |
|----------------------------|------------------------------------------|------------------------------------|
| Group A                    | Users arriving via <b>paid search</b>    | The same users, <b>with</b> ads    |
| Group B                    | Users arriving via <b>other channels</b> | <b>The same users, without ads</b> |
| The difference measures... | Ad effect + <b>intent difference</b>     | The ad effect. <b>Nothing else</b> |

The needed Group B is a **counterfactual** — a world we can never observe. All of causal inference is about **constructing a credible stand-in for it** — and how well a stand-in can be constructed depends on **how the data came into being**. This one idea organizes the entire course.

# 1.3 Types of research questions

# Every analytics task answers one of four questions

| # | Question                | Type         | Marketing examples                                                 |
|---|-------------------------|--------------|--------------------------------------------------------------------|
| 1 | What is going on?       | Descriptive  | Segment sizes, churn rate, funnel drop-off, satisfaction tracking  |
| 2 | What will happen?       | Predictive   | Which customers will churn next quarter? Demand next month?        |
| 3 | What happens if we act? | Causal       | Does the discount <i>cause</i> extra sales? Did the campaign work? |
| 4 | What should we do?      | Prescriptive | Set the price; allocate budget across channels                     |

**Prescription requires causation.** To choose an action you must know what each action would *cause* — i.e., compare counterfactuals. Prediction alone is not enough... and confusing 2 with 3 is precisely the eBay mistake.

# Four types of questions, four kinds of analytics

|                                   | Description                                                                      | Prediction                                                                           | Causal inference                                                                         | Prescription                                                          |
|-----------------------------------|----------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------|-----------------------------------------------------------------------|
| <b>Typical question</b>           | How can our customers be grouped into segments based on their characteristics?   | Which of last year's visitors will purchase from our store next month?               | Does behavioral targeting <i>increase</i> the probability of purchase — and by how much? | Which customers should we target, at what budget, to maximize profit? |
| <b>What the data must contain</b> | <i>Features:</i> user characteristics (age, location, ...), products viewed, ... | <i>Inputs</i> (features) and the <i>output</i> to be predicted (purchase next month) | <i>Outcome, treatment,</i> and — without an experiment — the relevant <i>confounders</i> | Causal effect estimates <i>plus</i> costs, margins, and constraints   |
| <b>Typical analytics</b>          | Cluster analysis, descriptive statistics, visualization                          | Regression, trees, random forests, neural networks                                   | Randomized experiments; regression, diff-in-diff, IV, RDD                                | Targeting & uplift models, optimization, simulation                   |

Based on the classification of data-science tasks in Hernán, Hsu & Healy (2019), *Chance* — prescription column added.

# Prediction $\neq$ causation: same data, different questions

“Customers who see our newsletter churn 40% less.”

As a predictive statement ✓

Newsletter readership *predicts* retention — useful for **targeting**: flag non-readers as churn risks.

*Valid even if the newsletter does nothing.*

As a causal statement ✗ (so far)


“Send more newsletters  $\rightarrow$  churn drops”?  
Only if readership *causes* retention — but engaged customers **both** read newsletters **and** stay.

*Requires a counterfactual comparison.*

**Classify these (2 min, neighbor discussion):** (a) “Which ad creative will get the highest CTR?” (b) “How large is our Gen-Z segment?” (c) “Should we raise the price by 5%?”


# 1.4 Correlation and causation

# Did Internet Explorer prevent murders?




A **spurious correlation**: two variables move together **without any causal link** — here simply because both **trend in the same direction over time**. Correlation is cheap; causation is not.

# Two pictures of “X relates to Y”



Same newsletter, same chart — **and the true effect is +3 in both worlds.** Panel (a) recovers it; panel (b) reports **+11**, because customers who subscribe differ from those who don't *before* any newsletter arrives. Which panel is which is **invisible in the plot**: it lies in **how the data came to be** (the *data-generating process*).

# What generates panel (b)



*observed in (b): +11 = +3 effect + ~8 selection*

The confounder is the customer's **underlying brand interest**: it makes people **subscribe** *and* it makes them **buy**. Randomizing delivery — panel (a) — cuts its left arrow, and only the **+3** remains (more on this later).


# Four ways to get the same scatter plot

# The four data-generating processes


1 · Causal




2 · Confounded



3 · Spurious



4 · Simpson



Four indistinguishable scatter plots — four different worlds. The plot shows the **correlation**; only the **DGP** says what X does to Y — and thus what happens if you **act** on X.

# Reverse causality: who is causing whom?

Observed: regions with higher ad spend show higher sales.



*Story 1:* advertising drives sales.

*Story 2:* budgets are **set as a % of sales** — sales drive ad spend.

Both stories produce identical data. Marketing budget rules make reverse causality a **routine** hazard in marketing data, not an exotic one.

# The classical conditions for causality

For X to be accepted as a **cause** of Y, three conditions must hold:

1. **Concomitant variation** — X and Y vary together as the hypothesis predicts. *The data establish this directly — all four scatter panels satisfy it. It is the cheap condition.*
2. **Temporal order** — the cause precedes the effect. *Sometimes checkable in data — but last slide's ad-budget example shows how it fails: sales also drive spending.*
3. **Absence of other possible causal factors** — no alternative explanation accounts for the association. *No dataset can certify this on its own: intent, segments, time, tiers...*

The entire difficulty of causal inference is **condition 3** — and this course teaches **strategies for making it plausible**.

Classical formulation after Field, Miles & Field (2012) — see the course book, §1.4.


# 1.5 Variables and their roles

# The two lead roles





- **Dependent variable (DV)** — the *outcome*: the quantity the manager cares about (purchases, churn, willingness to pay)
- **Independent variable (IV)** — the *presumed cause*: the variable whose influence we want to establish (discount, ad exposure, price)
- In **experiments**, the independent variable under the researcher's control is called the **treatment**; its values define the *experimental conditions* (treatment vs. control group)

# The villain: the confounder



- Causes **both** treatment and outcome
- Creates correlation **without** causation — panel 2 of our gallery
- **The eBay confounder:** purchase intent → people search for “eBay” → see the ad → and buy anyway
- Whole toolkits (Ch. 5–6) exist to deal with confounders in observational data

# Supporting roles: mediator and moderator



## Mediator — the *how*


Lies **on the path**: the mechanism through which X works.  
Why care? It tells you **what to optimize** (creative that builds awareness).

## Moderator — the *for whom / when*

Points **at the arrow**: changes the **strength** of the effect.  
“The discount works for new customers, not loyal ones.”


Today: **recognize** the roles. Chapter 5 teaches how to **handle** each one — and shows that mishandling them (e.g., “controlling” for a mediator) can change the conclusions from an analysis.

# The full cast around one causal question



# 1.6 Randomization

# Randomization cuts every arrow into the treatment




- Assignment depends **only on the coin** — not on intent, engagement, income, mood, ...
- **Systematic variation** — differences in the outcome created **by the treatment**: this is what we want to measure
- **Unsystematic variation** — differences from everything else (age, mood, time of day, ...): randomization spreads it **evenly across groups**
- So group differences = **treatment effect + noise**. No bias — and the noise is quantified in Module 3

Randomization does **not remove** other influences on the outcome — it makes them **unsystematic**: noise instead of bias. Noise averages out; bias does not.

# “Balanced — even on what we never measured”

**Covariate balance: exposed vs. unexposed group (simulated, n = 2,000)**

Standardized mean difference per variable · dashed lines: conventional  $\pm 0.1$  band



Left: the groups differ systematically — including on **intent**, which we (realistically) never observe. Right: one coin flip balances everything at once, **without knowing what needs balancing**. That is the magic randomization buys.

# Two ways to randomize: between vs. within subjects

|            | Between-subjects                           | Within-subjects                                                               |
|------------|--------------------------------------------|-------------------------------------------------------------------------------|
| Design     | Each person sees <b>one</b> condition      | Each person sees <b>all</b> conditions (repeated measures)                    |
| Example    | Group A: targeted ad · Group B: generic ad | Everyone rates ad no. 1, then ad no. 2                                        |
| Comparison | <b>Across people:</b> mean(A) vs. mean(B)  | <b>Within person:</b> condition 1 vs. condition 2                             |
| Strength   | No carryover between conditions            | Each person is their <b>own control</b> → far less noise, fewer participants  |
| Watch out  | Needs more participants                    | <b>Order &amp; carryover effects</b> → randomize the order (counterbalancing) |

Both are randomized experiments — they differ in **what is randomized**: *who gets which condition vs. the order of conditions*. Full treatment in Module 4.

# Why the industry runs on experiments

*“Our success at Amazon is a function of how many experiments we do per year, per month, per week, per day.”*

— Jeff Bezos

*“At any given point in time, there isn’t just one version of Facebook running, there are probably 10,000.”*

— Mark Zuckerberg

**A/B test** = industry name for a between-subjects randomized experiment, run continuously and at scale.

# Guidelines for field experiments

1. **Choose the unit of randomization.** Finer units (users rather than regions) → more statistical power — but weigh against logistics and spillovers. In digital environments, user-level A/B tests make fine granularity cheap
2. **Watch for spillovers & crossover.** Treated units can affect untreated ones (product adoption spreads to friends; a tipping prompt changes driver behavior for *all* riders — Uber, Chandar et al. 2019). If spillovers loom, randomize at a **higher level** (households, cities). *Crossover*: a unit meant for control gets treated anyway (multiple devices, shared accounts)
3. **Complete vs. stratified randomization.** Pure coin flips can leave small samples imbalanced by chance. **Stratify**: split units into homogeneous subgroups first, randomize within each — worth it when a covariate strongly predicts the outcome

Lambrecht & Tucker (2016); Chandar, Gneezy, List & Muir (2019).

# The price of observational data

## Experimental research

*“To find out what happens when you change something, it is necessary to change it.”*

— Box, Hunter & Hunter (1978)

## Observational research

*“Behind any causal conclusion there must be some causal assumption, untested in observational studies.”*

— Pearl (2009)

That untested assumption has a name: **no unobserved confounders** — every variable driving both treatment and outcome has been measured and correctly adjusted for.

- **Untestable in principle:** you cannot check balance on variables you did not observe
- Experiments **buy** the assumption *by design*; observational studies must **argue** for it — with domain knowledge, not statistics
- Not a counsel of despair: the three *Observational data* sessions build the toolkit for making that argument credible

# 1.7 Quasi-experiments

# Sometimes the world randomizes for you (almost)

A spectrum of designs:

|                        | Experiment                 | Quasi-experiment                          | Observational                      |
|------------------------|----------------------------|-------------------------------------------|------------------------------------|
| Who assigns treatment? | The researcher (coin flip) | External event / rule                     | The units themselves               |
| Example                | A/B test of a discount     | Platform launches in some countries first | Newsletter readers vs. non-readers |
| Key requirement        | Correct execution          | “As-good-as-random” timing/rule           | <b>No unobserved confounders</b>   |

- Policy changes, staggered rollouts, regional boundaries → treatment variation **not chosen by the units**
- Workhorse method: **difference-in-differences** — compare *changes* in treated vs. untreated markets
- Extremely common in practice (and in marketing research) → dedicated session in **Module 6 (panel data)**

# 1.8 From business problem to research design

# The translation template

Manager: “*I want to know whether our loyalty program is worth the money.*”

| Design element     | Your translation                                                               |
|--------------------|--------------------------------------------------------------------------------|
| Outcome (DV)       | Customer contribution margin, 12 months                                        |
| Treatment (IV)     | Loyalty program membership                                                     |
| Unit of analysis   | Customer                                                                       |
| Question type      | Causal (→ prescriptive decision)                                               |
| Likely confounders | Purchase frequency & engagement <i>before</i> joining                          |
| Feasible design?   | Randomize invitations? Staggered regional rollout? Observational + adjustment? |

A precise, falsifiable formulation (“membership increases 12-month margin”) beats a vague aspiration (“the program creates value”) — it dictates data, design, and analysis.

# Your turn (pairs, 10 min)

Translate one of these into the template:

1. “Should we invest more in influencer marketing?”
2. “Is our app driving retention, or do loyal customers just download the app?”
3. “Did last spring’s rebranding damage sales in Germany?”

For each: DV, IV, unit, question type, two plausible confounders, and the **best feasible design** (experiment / quasi-experiment / observational).

# 1.9 The analytics workflow

# Where AI helps — and where it cannot replace you

| Workflow step               | Nature of the work | AI assistant...                                            |
|-----------------------------|--------------------|------------------------------------------------------------|
| Formulate the question      | <b>Judgment</b>    | can challenge you — cannot know your business              |
| Choose design & method      | <b>Judgment</b>    | suggests options — <i>you</i> own the assumptions          |
| Write & debug R code        | Mechanical         | <b>excellent</b> — use it freely, including in this course |
| Interpret results           | <b>Judgment</b>    | drafts text — but cannot verify the causal claim           |
| Communicate to stakeholders | Mixed              | strong support, needs your framing                         |

**Course policy in one line:** use AI to *accelerate* your learning of the judgment steps, not to *skip* them. Assessments in this course are designed so that the judgment steps are what earns the points.

# 1.10 Course overview

# One map for the whole course

| Module                               | Sessions                 | Driving question                                       |
|--------------------------------------|--------------------------|--------------------------------------------------------|
| 1 · Questions, decisions & data      | 05.10.                   | Which question am I really asking?<br>( <i>today</i> ) |
| 2 · Describing data                  | 12.10.                   | What is going on?                                      |
| 3 · Uncertainty                      | 19.10.                   | How sure can I be?                                     |
| 4 · Learning from experiments        | 19.10. / 27.10.          | Did it work?                                           |
| 5 · Learning from observational data | 02.11. / 09.11. / 16.11. | What if we cannot randomize?                           |
| 6 · Panel data & quasi-experiments   | 30.11.                   | What if the world randomizes for us?                   |
| 7 · From estimate to decision        | 07.01.                   | What should we do?                                     |
| Group project coaching               | 23.11. / part of 07.01.  | Your own experiment, end to end                        |
| Final exam (in person)               | 15.12.                   | —                                                      |

**Group project preview:** design a survey **containing a real randomized experiment**, field it, analyze it. Everything from today — counterfactuals, randomization, confounders — becomes something you *build*, not just recognize.

# Learning check — can you spot the reasoning?

1. An analyst reports: “*App users have 2× the lifetime value of non-users — the app clearly pays off.*” Which question type is being **answered**, which is being **claimed**? Name the most plausible confounder.
2. A platform’s advertising dashboard compares purchases of users who clicked an ad vs. users who saw it and didn’t click. Is this an experiment? What role does “interest in the product” play?
3. Your CMO wants to know whether to extend a price promotion nationally after a pilot in Vienna. What design element makes the pilot informative — and what is the counterfactual?
4. For each: newsletter → (*opens* →) purchases; discount effect stronger for students — mediator or moderator?

# Next session

## Module 2 · Describing data — *What is going on?*

- Getting data into R, cleaning it, and not trusting it too early
- Summary statistics and the visualizations that earn their place
- **Before the session:** complete the R onboarding module (installation *or* Posit Cloud) — link on the course site. Any AI assistant is an endorsed study aid for this.

Course book chapter 1 (consolidation reading for today) is online — including all simulated examples from these slides, with their code.