Marketing Analytics

Chapter 5 · Learning from Observational Data

Nils Wlömert
Vienna University of Economics and Business (WU)

November 2026

Part 1 · Observational data I

Before we start

Paper quiz on Chapter 4 *(10 min)*

- Individual, closed book, counts toward participation
- Afterwards: a few leftovers from Chapter 4 and a look at the quiz questions

Assignment 2 is running: deadline Group A Sun, Nov 8 · Group B Wed, Nov 11; debrief at the start of next session.

Devices only when we code

- **Input blocks:** laptops closed, phones away
- **Coding block** (40 min, at the end): laptops open, alone or with your neighbor; results via Menti, then two tables present

> Today: correlation and the regression line → reading a regression → adjustment and exogeneity → categorical predictors → prediction vs. explanation → **one coding block with two cases** at the end.

Today’s session

Observational data I *(today)*

5.1 Smart Mix without randomization
5.2 Correlation, the line, and least squares
5.3 Reading a regression
5.4 Adjusting for confounders — and exogeneity
5.5 Categorical predictors
5.6 Prediction vs. explanation

Observational data II *(next session)*

5.7 Functional form and elasticities
5.8 Logistic regression
5.9 Diagnostics that change decisions
5.10 Difference-in-differences

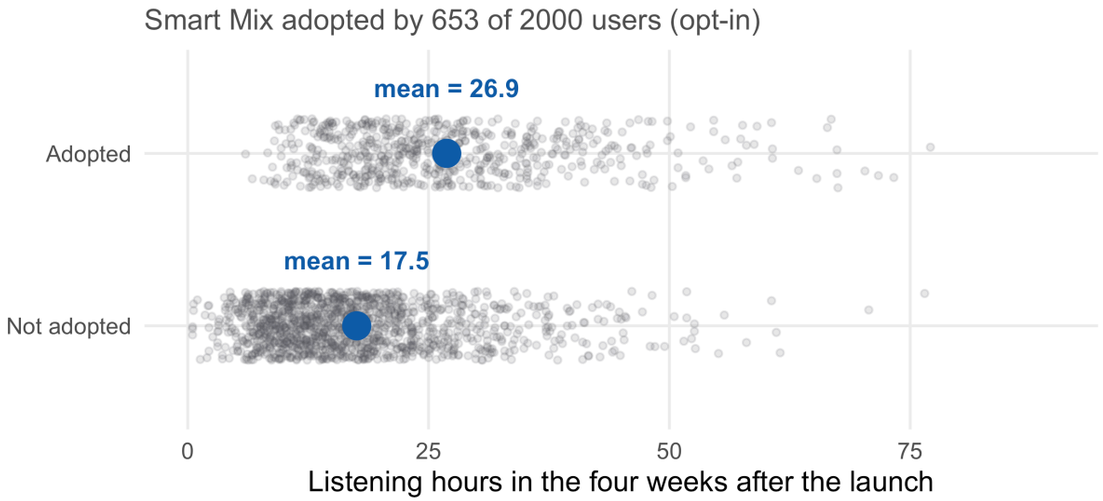
> **Driving question:** *What if we cannot randomize?*

5.1 Smart Mix without randomization

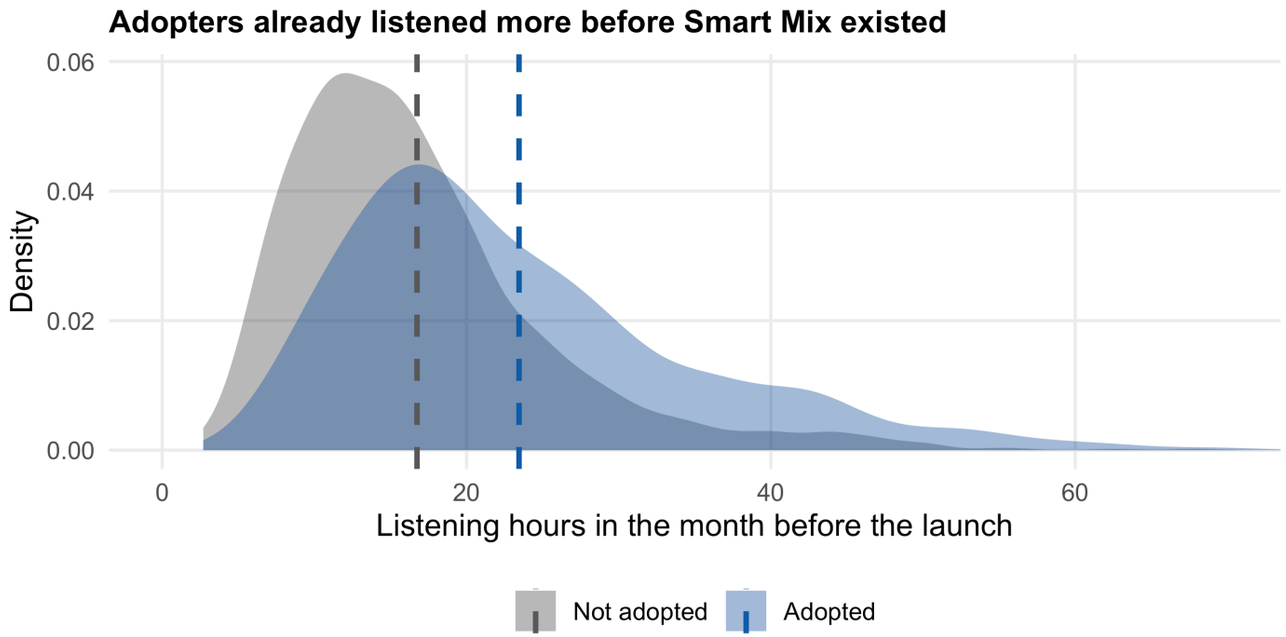
A different history

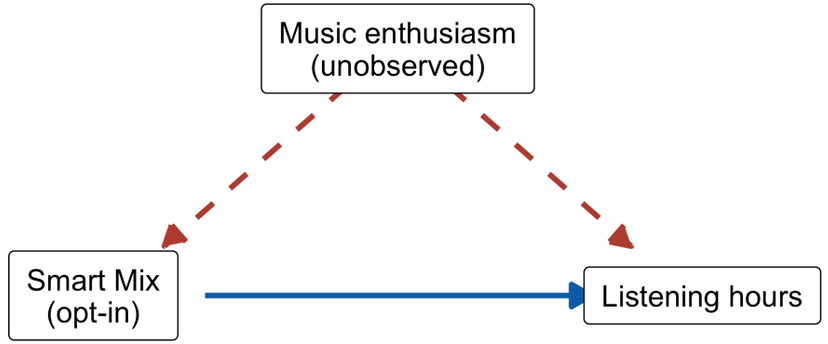
- Chapter 4: A/B test, 400 users → **+3.2 hours**
- Now: no experiment. Smart Mix launched as an **opt-in** feature
- Four weeks later: data of **2,000 users**, 653 adopters

The dashboard says: adopters listen **9.3 hours more** — three times the experimental effect. Would you report this number to management?



Adopters were different before Smart Mix existed

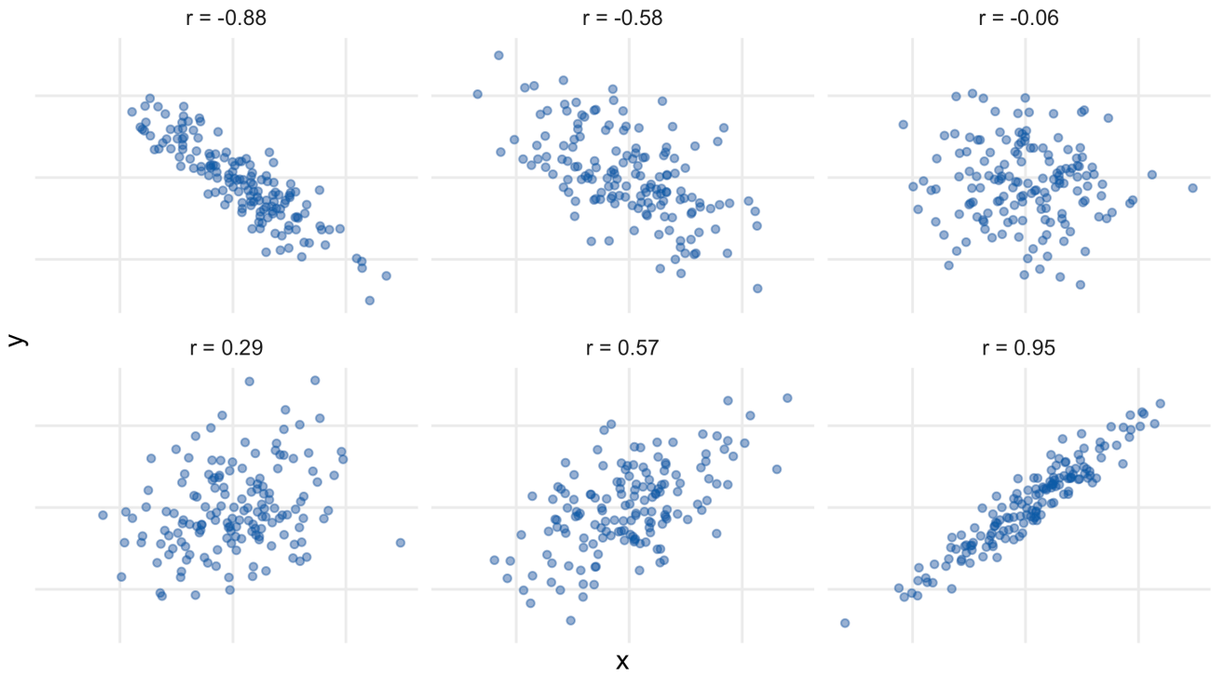




Same costume, new villain: **music enthusiasm** drives adoption *and* listening. Without a coin flip, we have one tool left: **adjustment**.

5.2 From correlation to a line

The correlation coefficient



$$r = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2}\sqrt{\sum_i (y_i - \bar{y})^2}}$$

- Standardized covariance → always between **−1 and +1**, no units
- **Sign** = direction · **size** = how tightly the cloud hugs a straight line
- Rules of thumb: 0.1 small · 0.3 medium · 0.5 large
- Smart Mix: `cor(prior_hours, hours)` = **0.64**

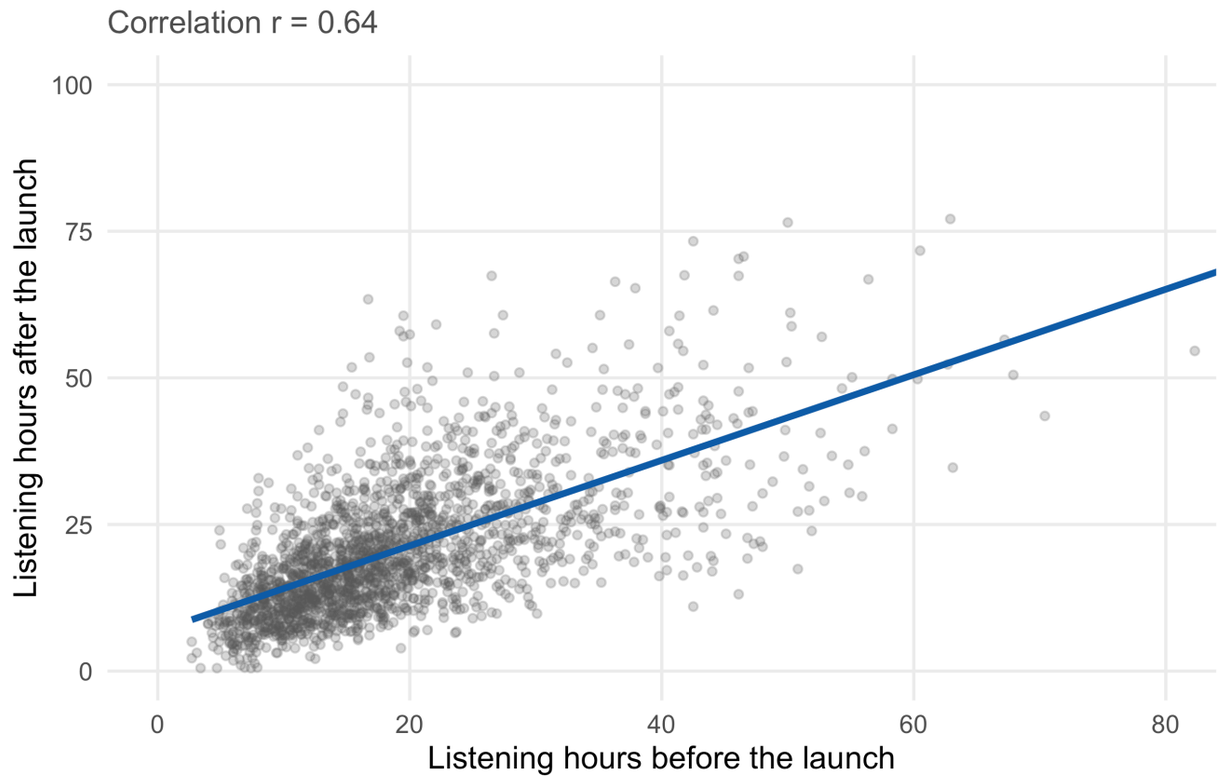
Two caveats: *r* measures **linear** association only (a U-shape can have $r \approx 0$) — and it is **symmetric**: nothing about which variable causes which.

Correlation vs. regression

	Correlation	Regression
Estimated coefficient	r, bounded between −1 and +1, no units	β_1, a priori not bounded, in units of y per unit of x
Interpretation	Linear association between two variables; **bidirectional** (r of x and y = r of y and x)	(Linear) relation of one or more predictors to an outcome; **directional** (x explains y)
Role of theory	Neither required nor testable	Required (which variable is the outcome, which predictors belong in the model) — and testable

> Correlation: *that* and *how tightly* two variables move together. Regression: *how much* y changes with x — which presupposes a direction, and that is where theory and this chapter's assumptions come in.

Two variables, one line



$$\text{hours}_i = \beta_0 + \beta_1 \text{ prior hours}_i + \varepsilon_i$$

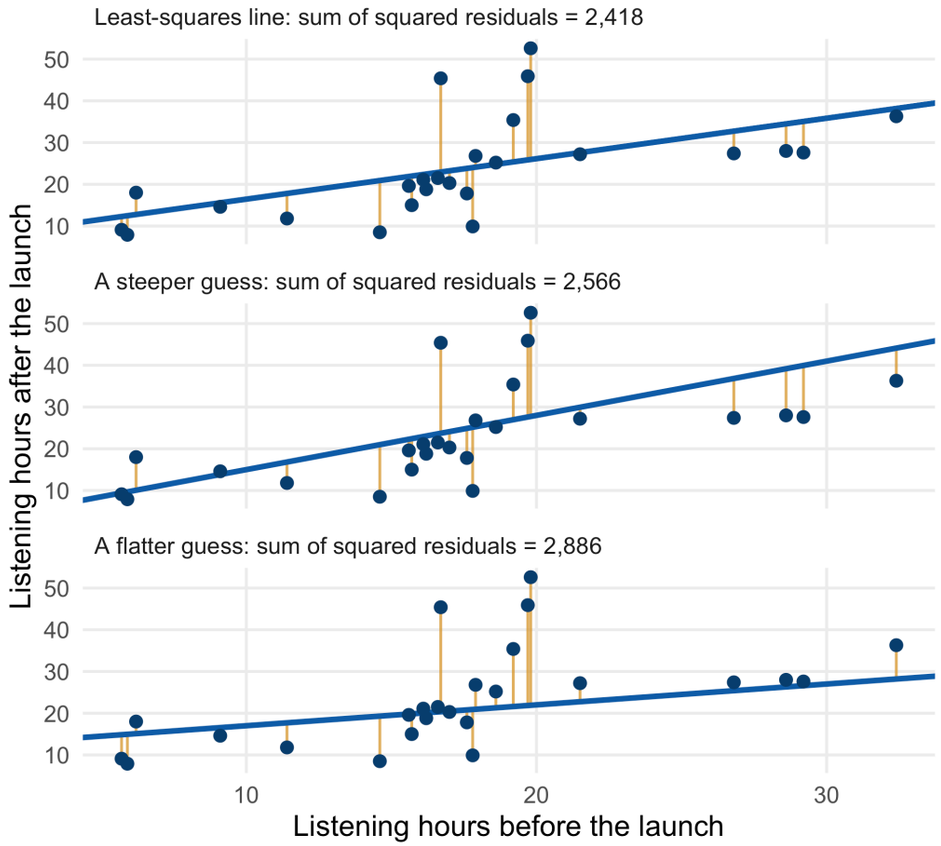
- i = one user · β_0 **intercept** · β_1 **slope** · ε_i **error term**
- Correlation: *that* they move together · regression: *how much*

```
1 coef(lm(hours ~ prior_hours, data = smartmix_obs))
```

```
(Intercept) prior_hours
     6.7849      0.7291
```

> Each extra hour before the launch **goes with** 0.73 hours after it — *goes with*, not *causes*.

Which line? Least squares



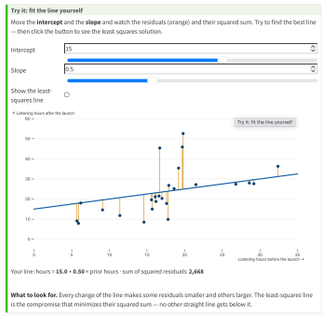
- Residual = observed − predicted (orange segments)
- **OLS** = the line that **minimizes the sum of squared residuals**:

[Math Processing Error]

- Squaring: big misses count a lot; above and below alike · `lm()` does exactly this

Live demo: fit the line yourself — course book, §5.2

https://wu-rds.github.io/MA2026/05-observational-data.html#sec-obs-line



5.3 Reading a regression

The full output

```r
model_naive <- lm(hours ~ smart_mix, data = smartmix_obs)
summary(model_naive)
```

```
Call:
lm(formula = hours ~ smart_mix, data = smartmix_obs)

Residuals:
   Min     1Q Median     3Q    Max
-20.88  -7.64  -2.06   4.76 101.22

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   17.541      0.300    58.5   <2e-16 ***
smart_mix      9.341      0.525    17.8   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 11 on 1998 degrees of freedom
Multiple R-squared:  0.137, Adjusted R-squared:  0.136
F-statistic:  316 on 1 and 1998 DF,  p-value: <2e-16
```

Reading the coefficients

`(Intercept)` = **17.5**: the predicted hours when all predictors are 0 — here the mean of the **non-adopters**.

`smart_mix` = **9.3**: the change in hours for a **one-unit increase** in the predictor. For a 0/1 variable that is the step from 0 to 1: **adopters − non-adopters**.

Continuous predictor (previous section): `prior_hours` = 0.73 → one more hour before the launch goes with **0.73 more hours** after it.

Several predictors (next section): every coefficient is read **"holding the other variables constant"**.

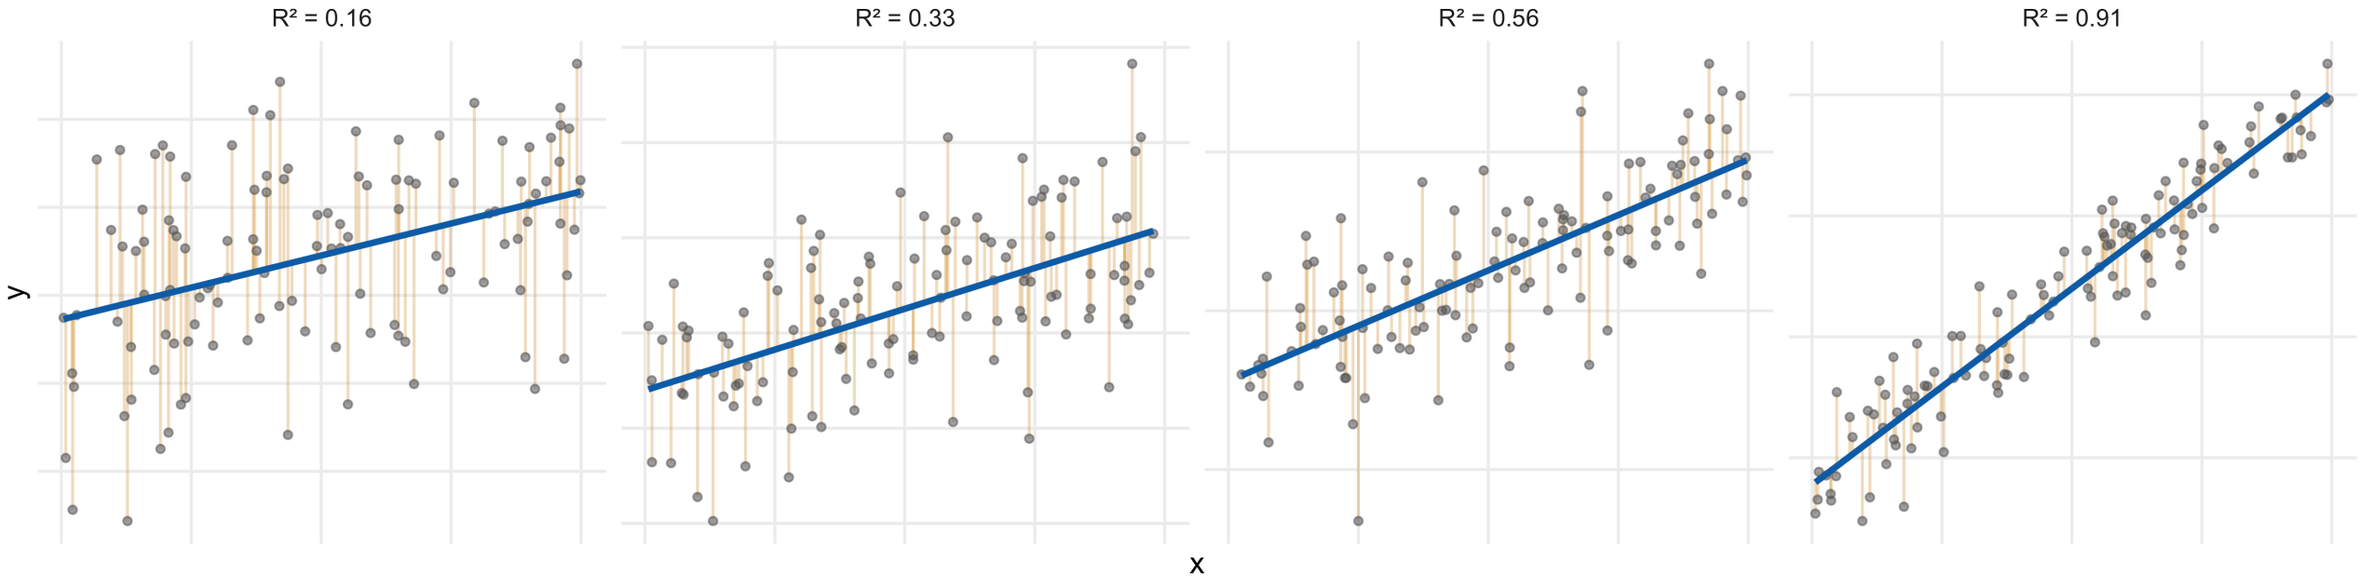
The general rule: a coefficient is the change in y when its predictor increases by **one unit** and nothing else changes. Whether that change is *causal* is the question of §5.4.

Line by line

Output	Meaning
`Estimate`	Intercept = mean of non-adopters (18.4); `smart_mix` = difference adopters − non-adopters (9.3)
`Std. Error`	Uncertainty of each coefficient (Chapter 3)
`t value`, `Pr(>|t|)`	*t*-test of each coefficient against zero (Chapter 4): signal / noise
`Residual standard error`	Typical distance of a user from the line, in hours
`Multiple R-squared`	Share of the variation in hours explained by the model (14%)
`F-statistic`	The whole model against "no predictors"

> With one 0/1 predictor, the regression **is** the comparison of two means — and its *t*-test is the two-sample *t*-test. `confint(model_naive)` gives the confidence intervals.

R^2: the share of explained variation



- Same idea as η^2 in the ANOVA: $R^2 = SS_M / SS_T$ — the share of the variation in y that the line explains
- Smart Mix dummy: 0.14 · prior listening: 0.41

> R^2 says how well the line **fits** — nothing about whether the slope is **causal**, and a small R^2 does not make an effect unimportant (14% explained, but 9 hours).

5.4 Adjusting for confounders

Holding prior listening constant

```r
model_prior <- lm(hours ~ smart_mix + prior_hours, data = smartmix_obs)
round(coef(summary(model_prior)), 3)
```

```
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    6.451      0.412   15.65        0
smart_mix      4.901      0.443   11.07        0
prior_hours    0.662      0.020   33.17        0
```

> The Smart Mix coefficient drops from **9.3** to **4.9**: now the difference between adopters and non-adopters **with the same prior listening**. Part of the naive gap belonged to prior listening — and is now attributed to it.

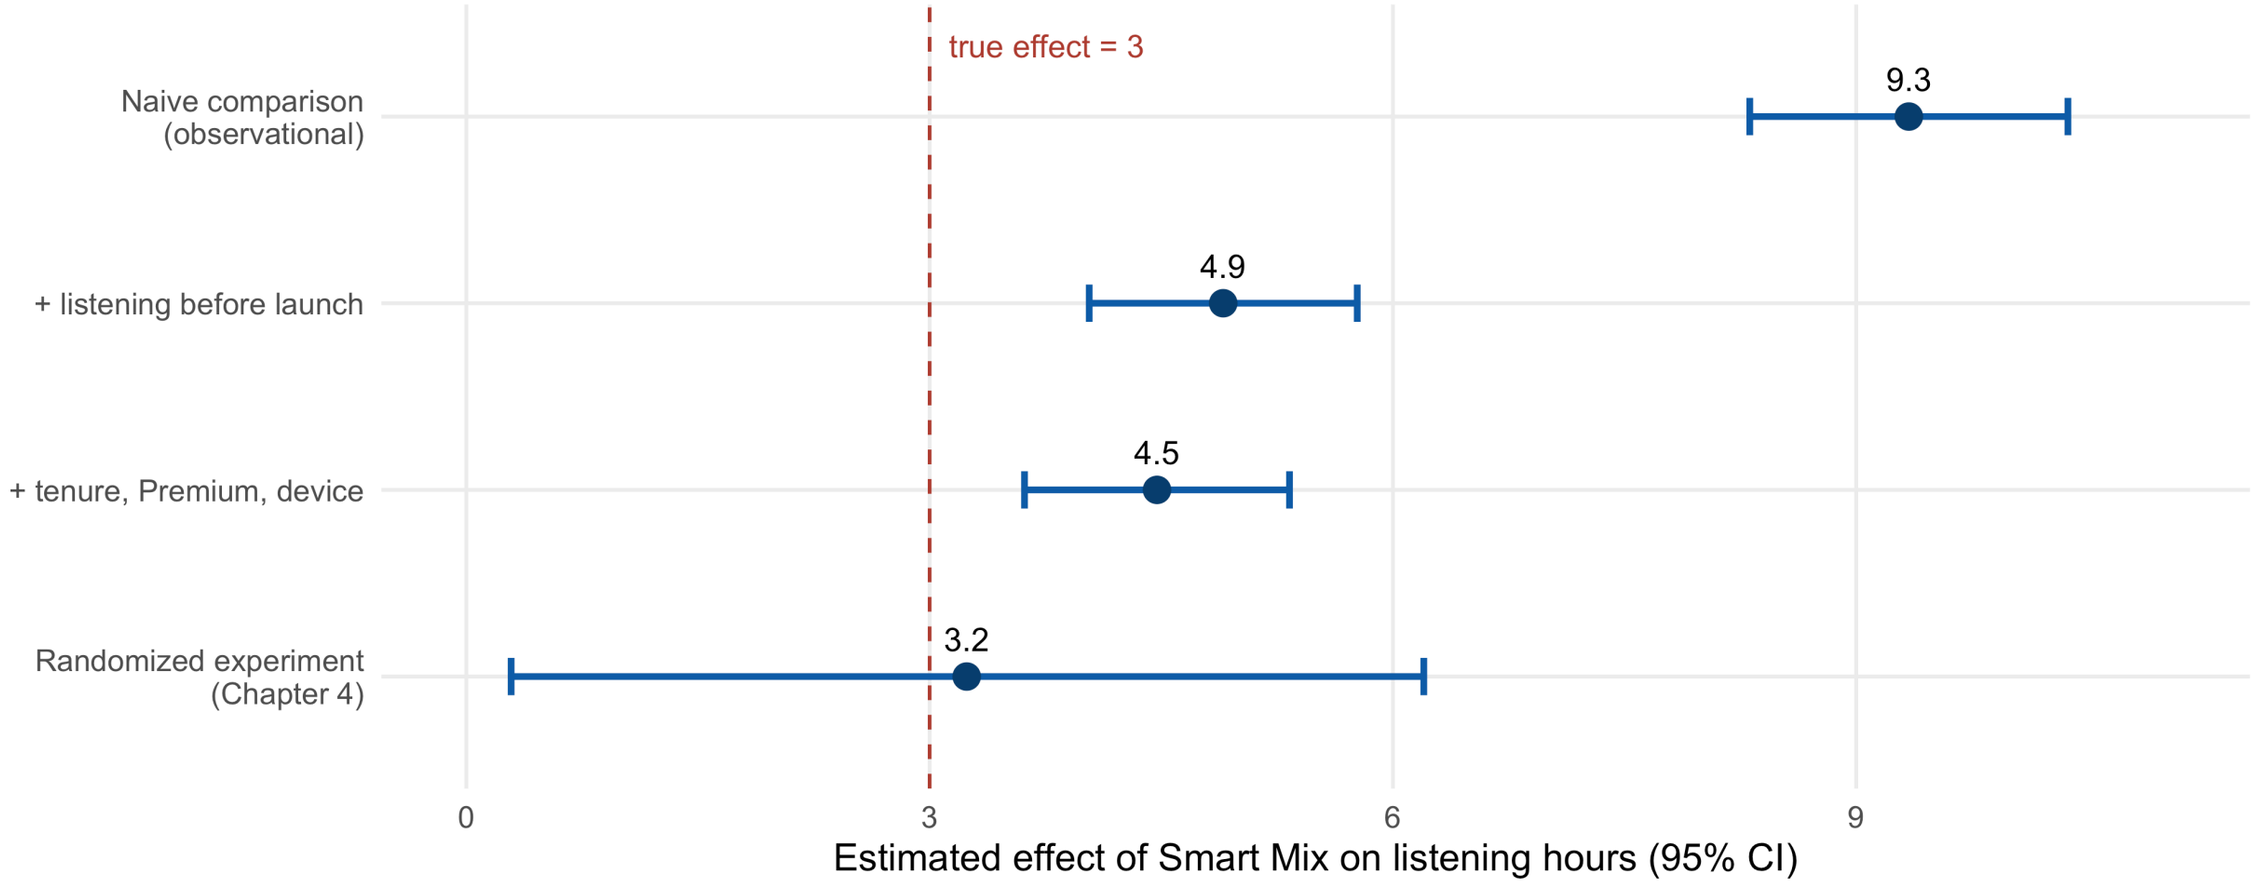
More variables, same logic

```
model_full <- lm(hours ~ smart_mix + prior_hours + tenure + premium + device, data = smartmix_obs)
round(coef(summary(model_full)), 3)
```

	Estimate	Std. Error	t value	Pr(>|t|)
(Intercept)	6.255	0.540	11.575	0.000
smart_mix	4.472	0.437	10.223	0.000
prior_hours	0.613	0.021	29.688	0.000
tenure	0.038	0.011	3.344	0.001
premium	2.924	0.426	6.869	0.000
deviceAndroid	−1.475	0.422	−3.493	0.000
deviceDesktop	−2.488	0.576	−4.320	0.000

Reading the table, one line each, "holding the other variables constant": `smart_mix` 4.5 = adopters listen 4.5 h more than non-adopters with the same prior listening, tenure, plan and device · `prior_hours` 0.61 = one more hour before the launch goes with 0.61 h more after it · `tenure` 0.04 = each month of tenure, +0.04 h · `premium` 2.9 = Premium users listen 2.9 h more than free users · `deviceAndroid` / `deviceDesktop` = difference to **iOS** (the reference). Tenure, Premium and device change little for Smart Mix: ≈ **4.5 hours**.

Naive, adjusted — and the truth

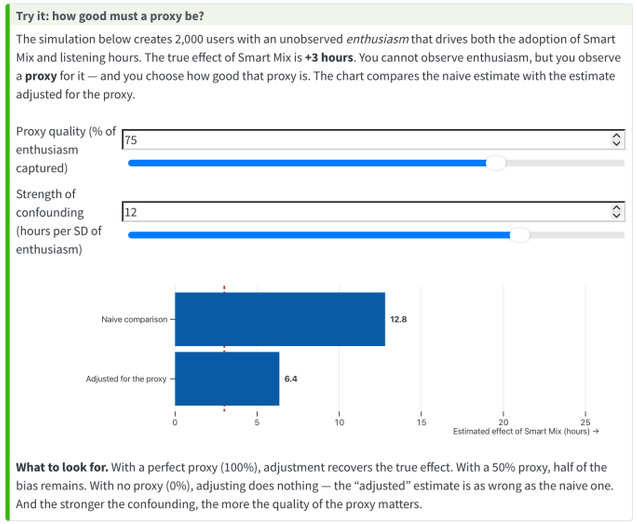


Adjustment removed ~75% of the bias — **not all of it**. We adjusted for a *proxy* of enthusiasm, not for enthusiasm. Regression removes exactly the confounding your variables capture.

Live demo: how good must a proxy be?

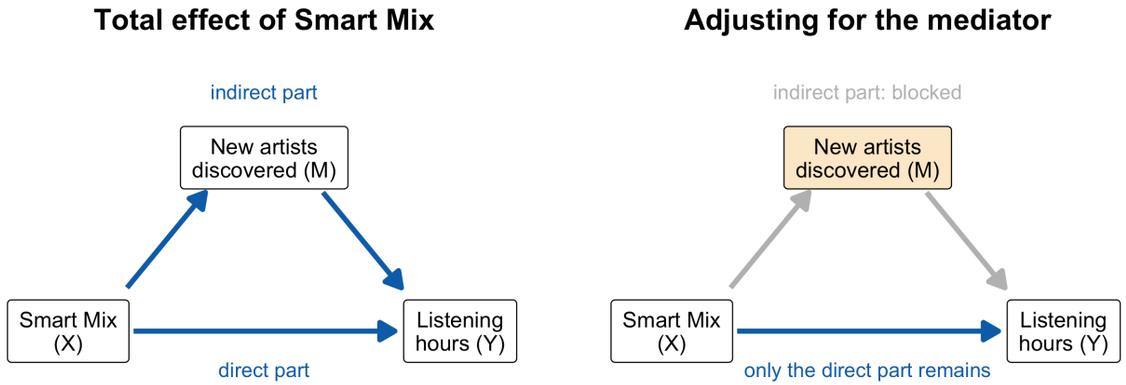
Course book, §5.4 — set the **quality of the proxy** and the **strength of confounding**.

https://wu-rds.github.io/MA2026/05-observational-data.html#sec-obs-adjust



> With a 50% proxy, how much of the bias is left? And what happens at 0%?

Two warnings



"Enough" is untestable. The assumption of no unobserved confounders cannot be checked in the data. Report *what* you adjusted for — and what you plausibly could not.

Adjusting for the wrong variable makes it worse. A **mediator** lies on the causal path: if Smart Mix works by making users discover new artists, adjusting for `new_artists` **blocks the indirect part** and the coefficient shows only the direct part. Only **confounders** — determined *before* the treatment, affecting treatment *and* outcome — belong in the model.

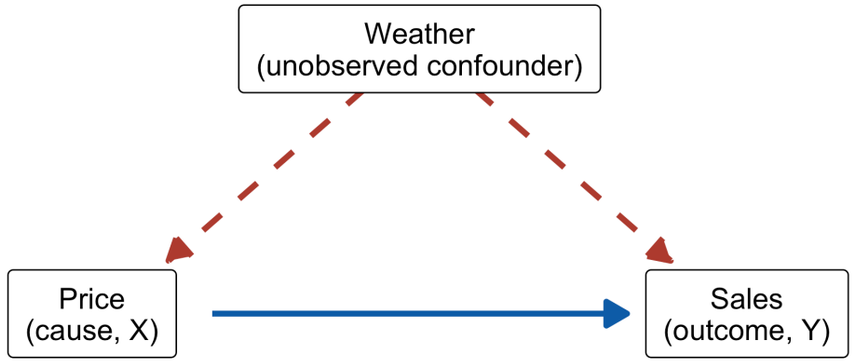
Where the unobserved variable goes

The true model (price and sales; weather raises both):

$$\text{Sales}_i = \beta_0 + \beta_1 \text{Price}_i + \beta_2 \text{Weather}_i + \phi_i$$

Weather is not in the data — so we estimate:

$$\text{Sales}_i = \beta_0 + \beta_1 \text{Price}_i + \underbrace{\varepsilon_i}_{= \beta_2 \text{Weather}_i + \phi_i}$$



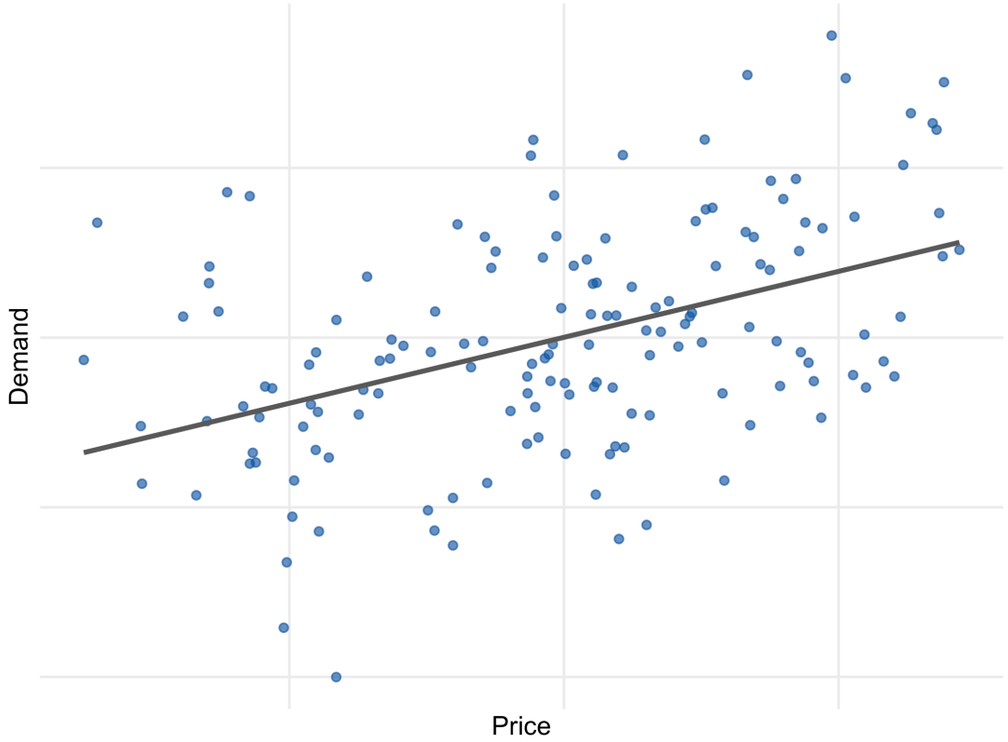
- The omitted variable does not disappear — it moves **into the error term**
- Weather also drives price → the error term is **correlated with the regressor**
- OLS attributes the weather effect to price: $\hat{\beta_1}$ is **biased**, and more data do not help

Smart Mix: replace price by **adoption**, weather by **enthusiasm**. Enthusiasm is in ε_i and drives adoption → 9.3 h instead of 3.

Variation has two components

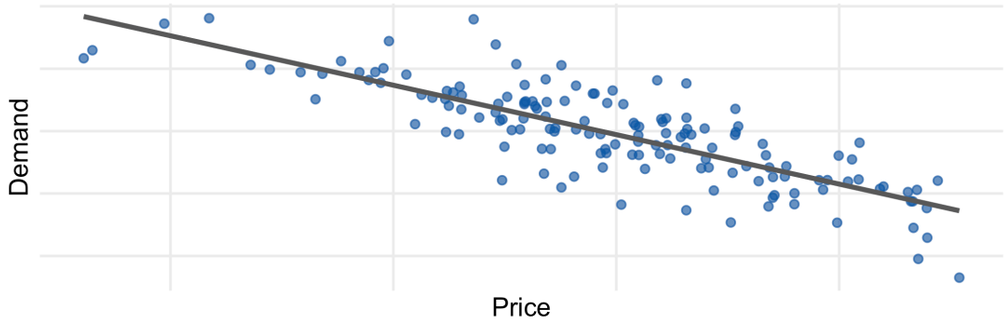
Observed variation

The data: higher prices, higher demand?



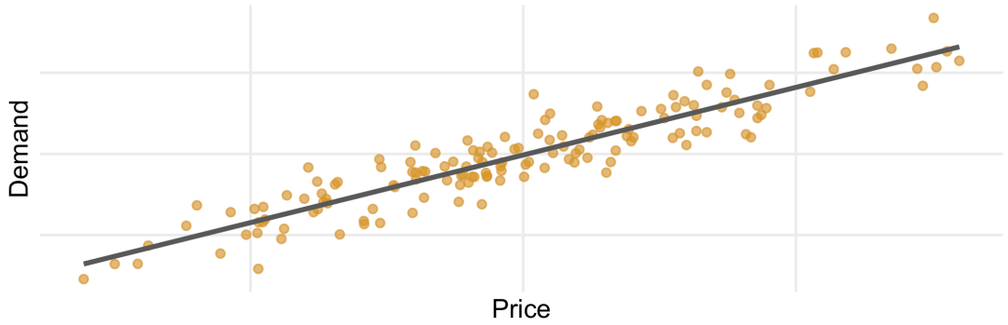
Exogenous variation

Price moves for its own reasons: the causal effect



Endogenous variation

Weather moves price AND demand: not the effect



The observed price variation mixes **exogenous** variation (price moves for its own reasons → the causal effect, downward) with **endogenous** variation (weather moves price *and* demand → upward). OLS on the mix gives a slope that is neither. Adjustment and experiments are ways to isolate the exogenous part.

Exogeneity

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

- The most important assumption when the goal is the **effect** of x on y
- The error term must not be related to x: $\text{Cov}(x_i, \varepsilon_i) = 0$ — x is **exogenous**
- If $\text{Cov}(x_i, \varepsilon_i) \neq 0$: x is **endogenous**, the estimate is **biased** (and stays biased with more data), and β_1 cannot be read as causal
- Randomization **guarantees** exogeneity. Opt-in **violates** it. Each adjustment variable moves a piece of ε_i into the model: `prior_hours` took most of the enthusiasm out (9.3 → 4.5); what is left still correlates with adoption (4.5 > 3)

Exogeneity cannot be tested. The residuals of a fitted regression are uncorrelated with the predictors *by construction* — the data always "look" exogenous. The argument has to come from outside: theory, and how the data came to be (who chose the treatment, and why).

Your turn (5 min)

Add **age** to the full model. Does the Smart Mix coefficient change? Why (not)?

```
model_age <- lm(hours ~ smart_mix + prior_hours + tenure + premium + device + age,
                data = smartmix_obs)
round(coef(summary(model_age)), 3)
```

Then: which variable in the data set should you **never** adjust for when estimating the effect of Smart Mix on listening — and why?

5.5 Categorical predictors

Dummies and the reference category

```
coef(model_full)[c("premium", "deviceAndroid", "deviceDesktop")]
```

```
premium deviceAndroid deviceDesktop
  2.924        −1.475        −2.488
```

- **0/1 variable** (`premium`): coefficient = difference between the two groups, other variables held constant
- **Several categories** (`device`): R picks a **reference** (first level: iOS) and creates one **dummy** per other category
- Each dummy = difference to the **reference**

Android users: **1.5 hours less** than iOS users; Desktop: **2.5 hours less** — holding everything else constant. *p*-values compare each category with the reference; “does device matter at all?” → `anova(model_without, model_full)` (the ANOVA of Chapter 4).

5.6 Prediction vs. explanation

Predicting listening time for new users

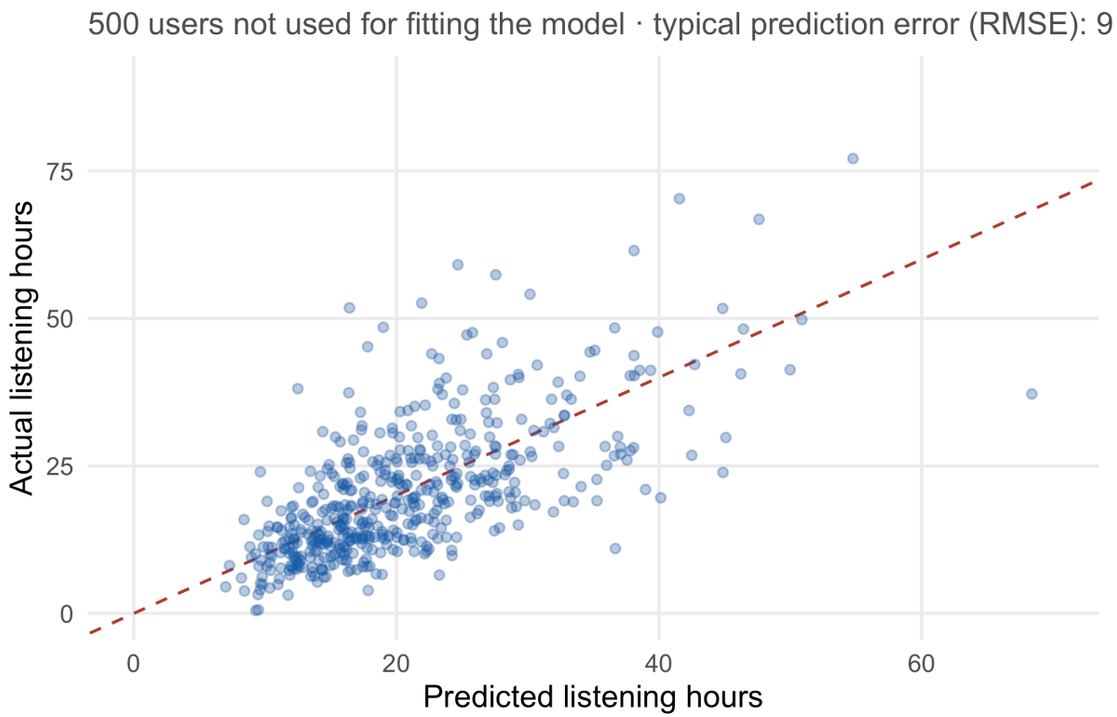
```r
set.seed(42)
train_ids <- sample(nrow(smartmix_obs), 1500)
train <- smartmix_obs[train_ids, ]
test  <- smartmix_obs[-train_ids, ]
model_pred <- lm(hours ~ prior_hours + tenure + premium +
                   device + smart_mix, data = train)
test$predicted <- predict(model_pred, newdata = test)

summary(model_pred)$r.squared        # in-sample R² (training data)
```

```
[1] 0.4817
```

```r
cor(test$predicted, test$hours)^2    # out-of-sample R² (test data)
```

```
[1] 0.4235
```



Fit on the training users, judge on the test users. The in-sample R^2 flatters the model (it was fitted to these very users); the **out-of-sample R^2** — the squared correlation between predicted and actual values in the test set — is the honest number. Close here (0.48 vs. 0.42): few predictors, many users. With many predictors and small samples, the gap is **overfitting** made visible.

How far off is a prediction? The RMSE

[Math Processing Error]

1. Each test user's **prediction error** $y_i - \hat{y_i}$
2. **Square** it: over- and under-predictions do not cancel, big misses count more
3. **Average** the squares
4. **Root**: back to the original unit (hours)

```
1  sqrt(mean((test$hours - test$predicted)^2))
```

```
[1] 9.425
```

> The RMSE is the **standard deviation of the prediction errors**: a typical prediction for a new user misses by about **9.5 hours**. For a decision maker, that is more useful than a share of explained variation — it is in the unit of the decision.

Two questions, two rules

Explanation (causal question): include confounders, exclude mediators, state what is unobserved. The **coefficient** matters; R^2 is secondary.

Prediction: include whatever predicts; judge by **out-of-sample errors**. Coefficients are not interpreted; the **predictions** matter.

A model can be a **good predictor and a bad explainer** at the same time. Where in your company would each mistake — reading causal effects off a prediction model, or judging a causal model by its R^2 — cost money?

Coding block: two cases (40 min)

Case 1: podcast promotion — `podcast_promo.csv`, 400 episodes: promotion spend, host followers, downloads.

1. Scatter plot of downloads against spend; the **correlation**; one sentence
2. `lm()` + `summary()`: interpret the **slope**; read off **R^2**
3. Is the slope **causal**? What would you need to argue?

Case 2: a loyalty program — `loyalty_program.csv`, 1,500 customers, membership is **opt-in**.

1. **Naive:** `spending ~ member` — interpret; why not causal?
2. **Adjust** for `spend_prev`, `tenure`, `device`: how does the coefficient change, what does it mean now? With 95% CI
3. Interpret the coefficient of `device`
4. Is membership **exogenous** after adjustment? Argue from the data-generating process

Laptops open — alone or with your neighbor. Section "Coding block" in `05-observational-data.R`.

> At the end: one table presents case 1, another table case 2 — Menti first: enter your slope (case 1) and your adjusted coefficient (case 2).

Learning check (Part I)

1. Slope of hours on prior hours = 0.73. Interpret — and why is it not causal?
2. Adding a variable can increase, decrease, or leave the Smart Mix coefficient unchanged. When does each happen?
3. `planFamily` = 4.2 with reference “Free”. Interpret. How do you test whether the plan matters at all?
4. R^2 = 0.71 for predicting next month’s spending with 25 predictors on 300 customers. What do you ask?

Next session

Observational data II — beyond the straight line, and a quasi-experiment

- Functional form: when a line is wrong — **elasticities**
- Yes/no outcomes: logistic regression
- The three diagnostics that change decisions
- **Difference-in-differences** with real playlist data

> Before next session: work through §5.1–5.6 on the course site. **Debrief of Assignment 2** at the start of the session.

Part 2 · Observational data II

Before we start

Debrief A2 + quiz Chapter 5, Part I *(15 min, one sheet)*

- Page 1: debrief of **Assignment 2** (3 assignment points)
- Page 2: quiz on **Chapter 5, §5.1–5.6** (participation)
- Closed book, no laptop

Assignment 3 is handed out today (end of session).

Devices only when we code — as last time

> Today: functional form and elasticities → logistic regression → diagnostics → difference-in-differences → **one coding block with two cases** → A3.

5.7 Functional form: when the line is wrong

Two assumptions of the linear model — and what relaxes them

$$y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \cdots + \varepsilon_i$$

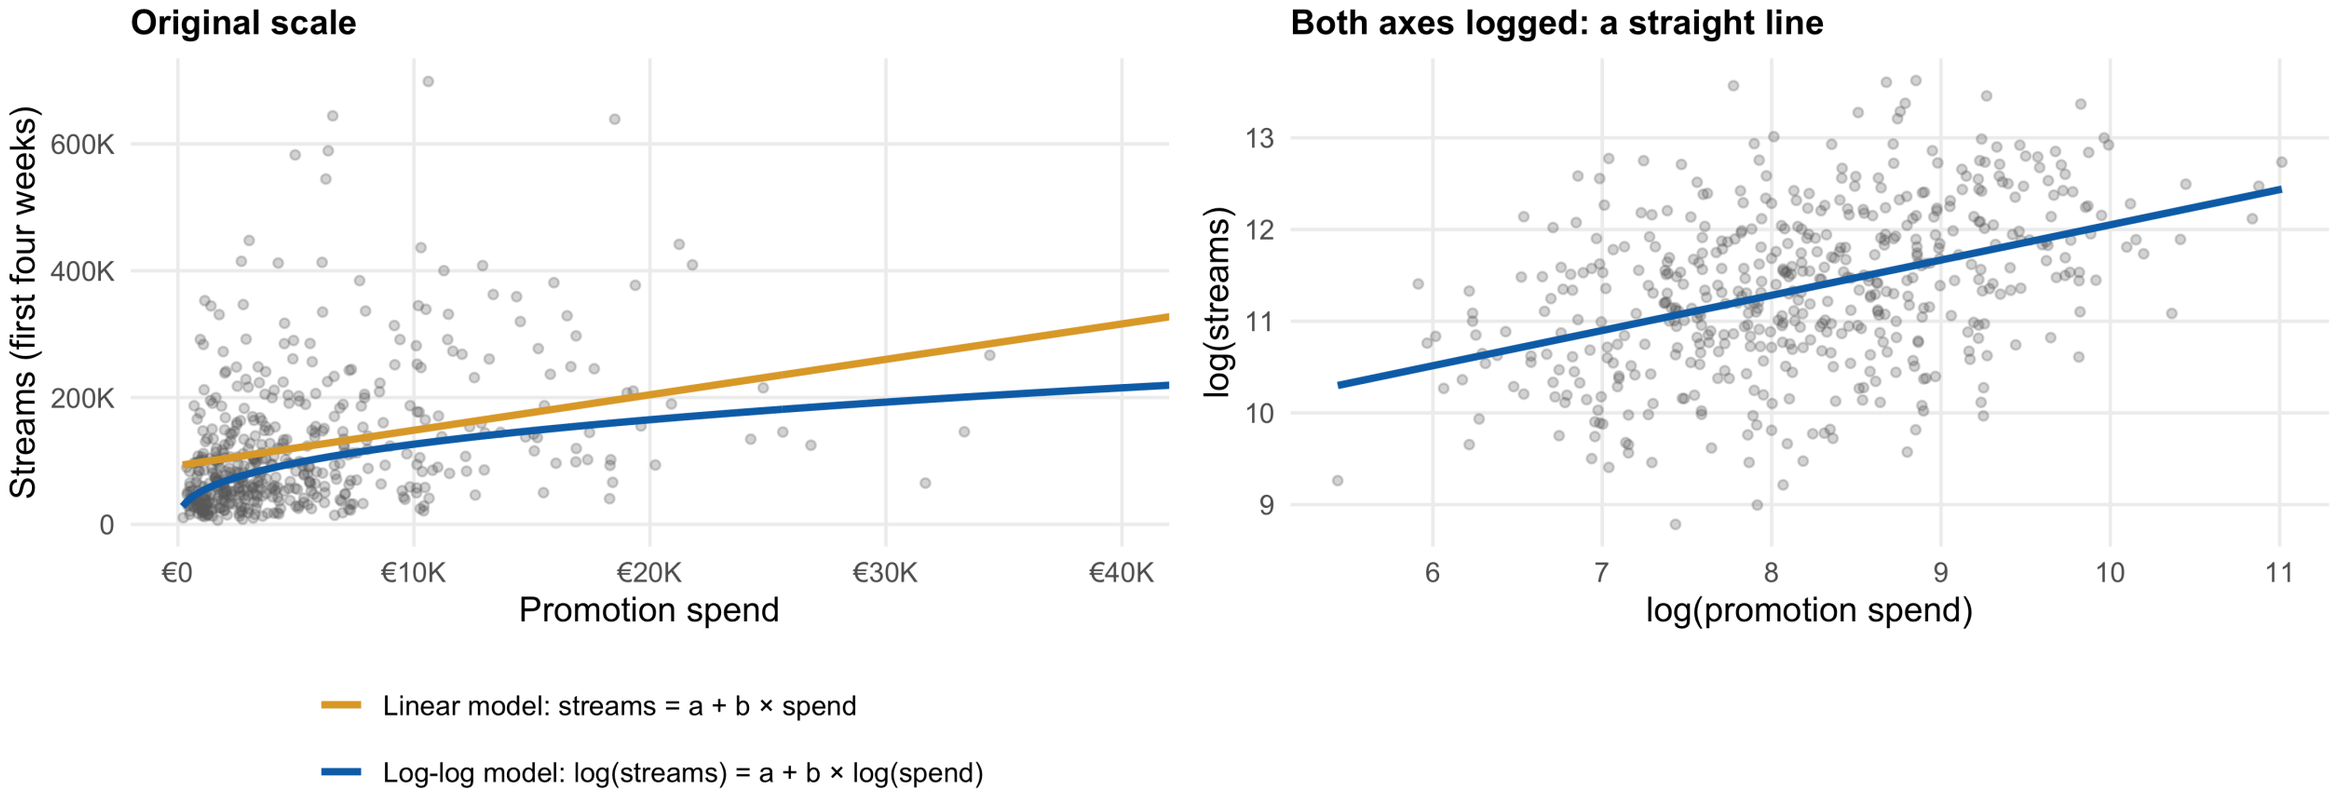
Linearity: a one-unit change in x has the same effect at every level of x — **constant marginal returns**. Ad spend 10 → 11 € adds as much as 100,000 → 100,001 €. No saturation.

Additivity: the effect of one predictor does not depend on the others. Each playlist placement adds the same number of streams for an artist with 1,000 or a million followers.

Where this course relaxes them:

- Linearity → the **multiplicative (log-log) model**, this section: diminishing returns built in
- Additivity → the multiplicative model (effects scale each other) and **interaction terms**, Chapter 6
- **Detecting** a wrong form → the residual plot, §5.9: curved residuals mean the straight line is wrong

Diminishing returns



500 promotion campaigns: the first euros buy many streams, later euros fewer. A straight line cannot bend — on **log scales**, the cloud becomes a line again.

The multiplicative model — and why logs make it linear

Multiplicative: each predictor *scales* the outcome

$$\text{streams}_i = e^{\beta_0} \cdot \text{spend}_i^{\beta_1} \cdot \text{followers}_i^{\beta_2} \cdot e^{\varepsilon_i}$$

Take logs on both sides → additive again, a job for `lm()`:

$$\log(\text{streams}_i) = \beta_0 + \beta_1 \log(\text{spend}_i) + \beta_2 \log(\text{followers}_i) + \varepsilon_i$$

Three log rules do the work:

$\log(a \cdot b) = \log a + \log b$ $\log(a^b) = b \log a$

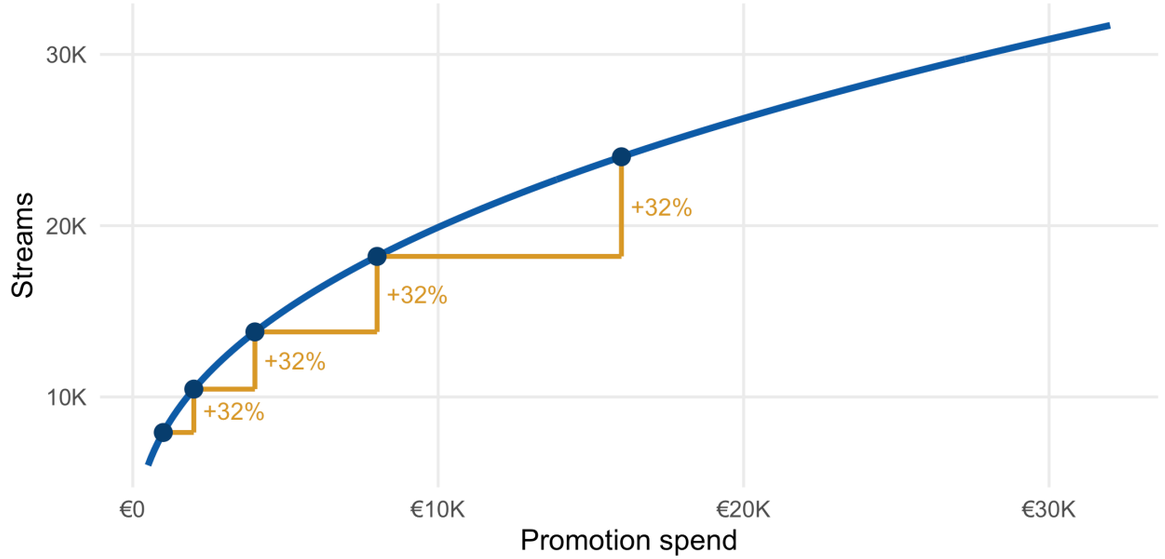
$\log(e^x) = x$

Additive model: each euro adds β_1 streams, always.
Multiplicative model: each euro adds a **percentage** of the current streams — diminishing returns built in.

Why the log transformation straightens the curve

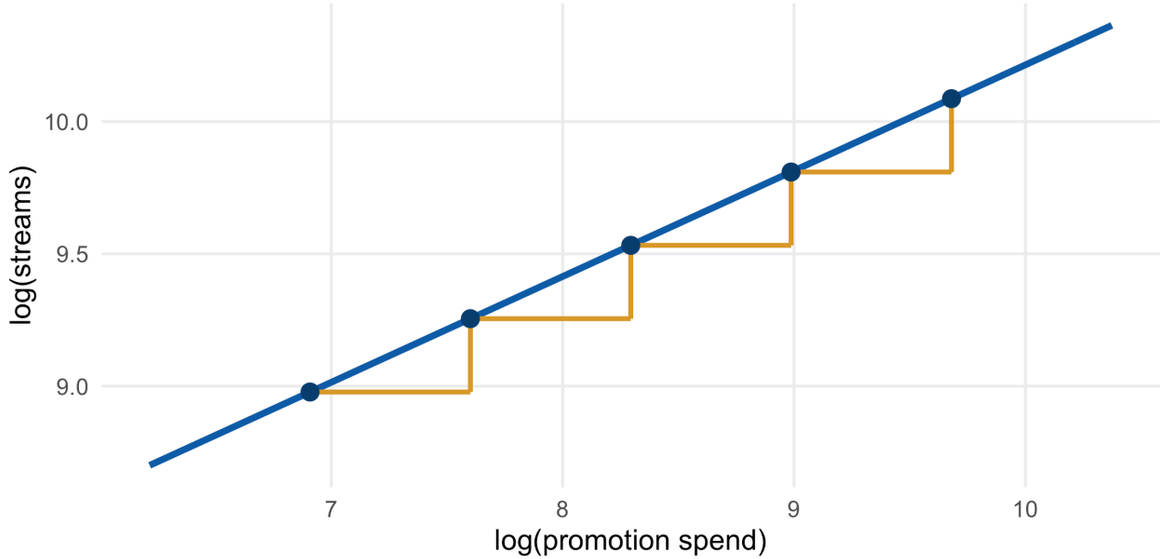
Original scale

Each doubling of spend (+100%) adds the same +32% to streams, but a shrinking number of streams



Both axes logged

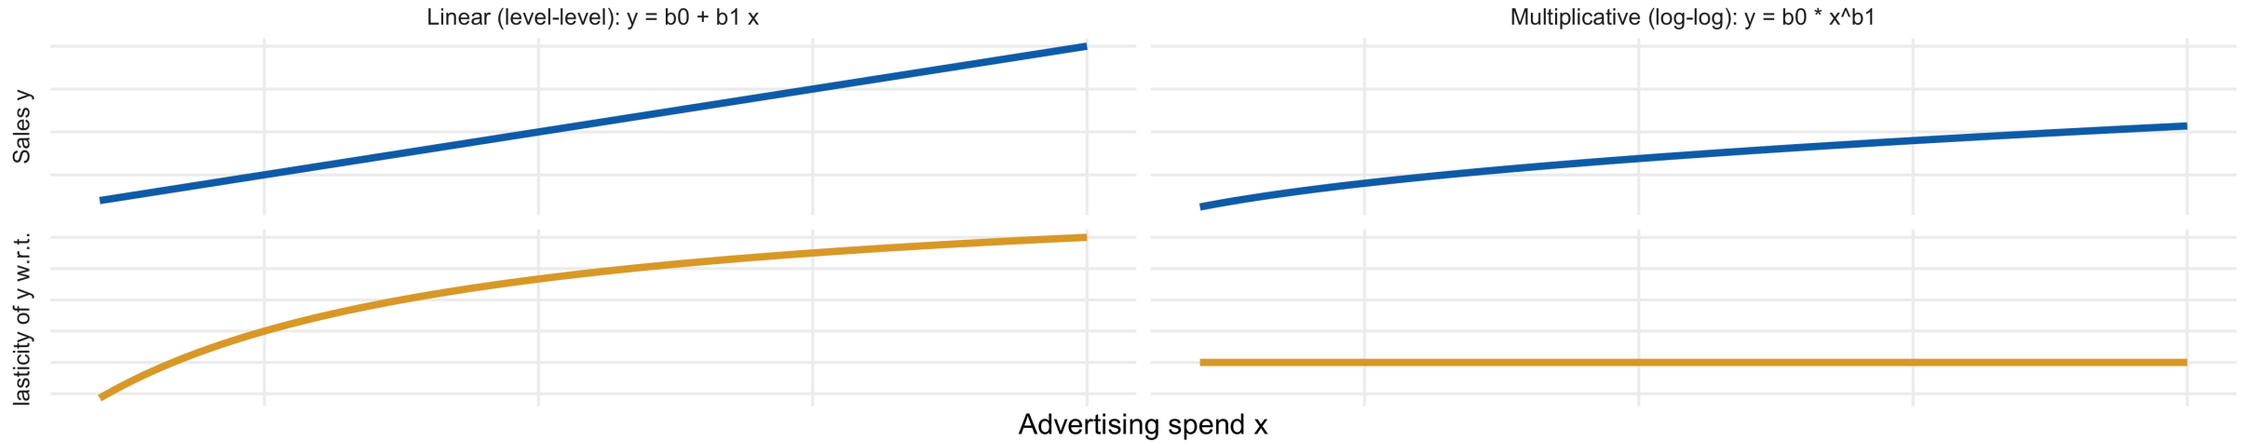
Equal percentage steps become equal distances: the curve is a straight line with slope 0.4



Multiplicative model: equal **percentage** steps in x → equal **percentage** steps in y (doubling spend always adds +32% streams, but ever fewer streams). The logarithm turns percentages into distances — a doubling is always $\log 2$ — so the curve becomes a line with slope = elasticity.

Elasticity: linear vs. log-log model

Linear: constant marginal returns, rising elasticity. Multiplicative: diminishing returns, constant elasticity.

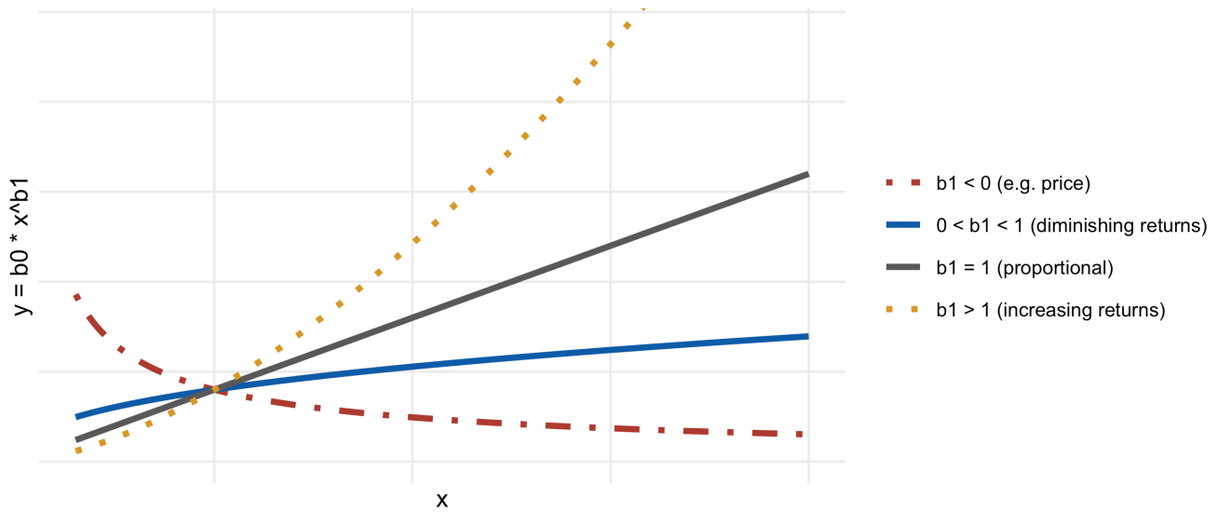


Elasticity $\eta_x = (\Delta y/y)/(\Delta x/x)$: % change in y per % change in x

- **Linear:** $\eta_x = \beta_1 \; x/y$ **rises** with x — implausible
- **Multiplicative:** $\eta_x = \beta_1$ **constant**; marginal returns fall

> Additive in **logs** = multiplicative in **levels**: the same *percentage* effect for every artist; the *number* of streams scales with the follower base.

What the elasticity says about the shape



- $\beta_1 < 0$: falls, flattening out — **price**
- $0 < \beta_1 < 1$: **diminishing returns** — advertising, promotion; most marketing elasticities (usually < 0.5)
- $\beta_1 = 1$: **proportional** — a straight line through the origin; the one case where the mean response is linear (no intercept, errors still multiplicative)
- $\beta_1 > 1$: **increasing returns** — rare; network effects

> Not defined at $x = 0$, no minimum level of y. Linear model: optimum at zero (price) or infinity (advertising); multiplicative: a finite optimum from elasticities and costs.

The log-log model: elasticities

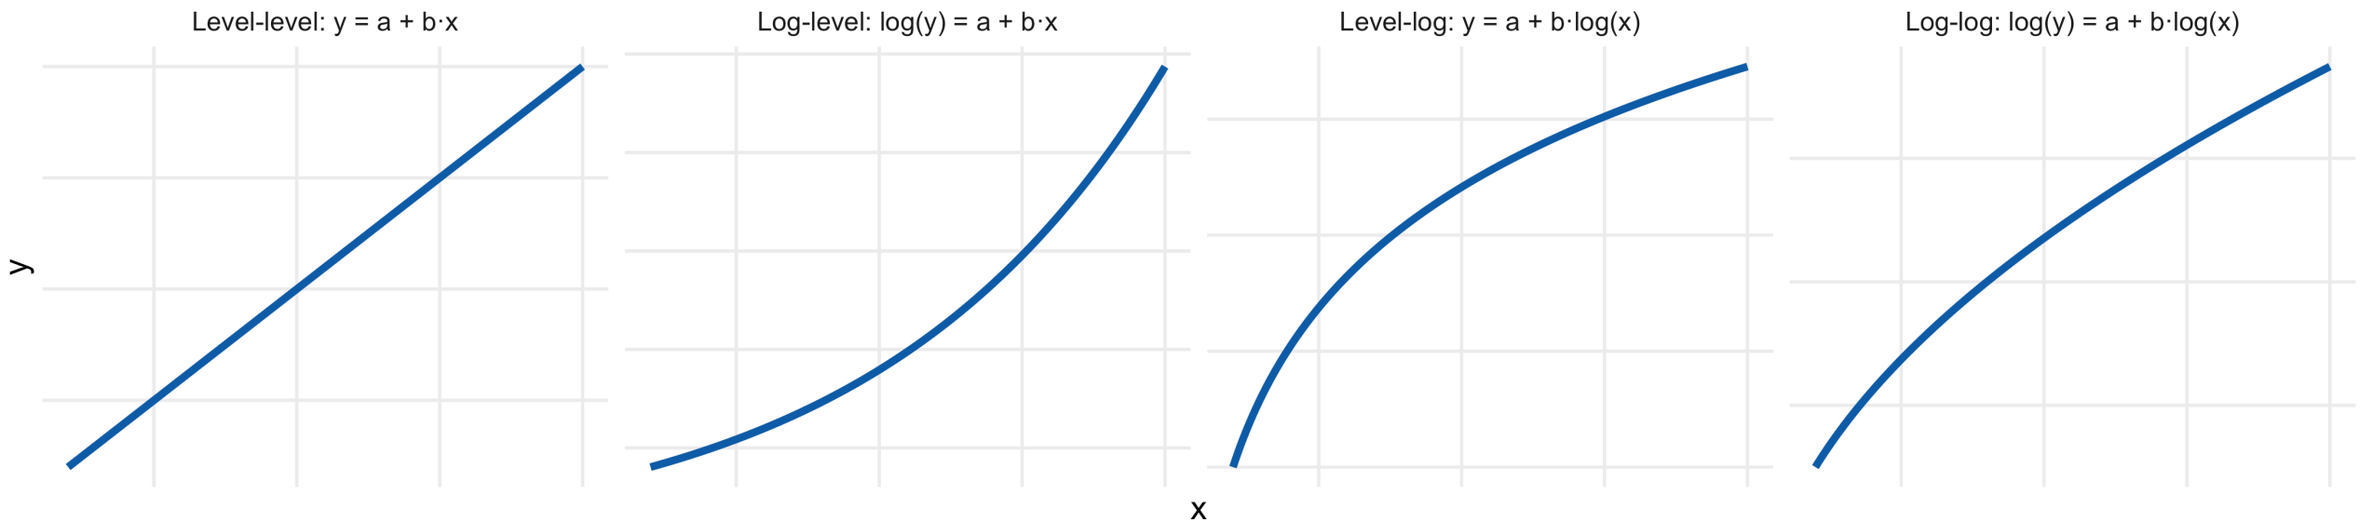
```
model_loglog <- lm(log(streams) ~ log(spend) + log(followers) + placements + genre, data = campaigns)
round(coef(summary(model_loglog)), 3)
```

```
                  Estimate Std. Error t value Pr(>|t|)
(Intercept)          2.881      0.235  12.282    0.000
log(spend)           0.331      0.021  16.001    0.000
log(followers)       0.507      0.016  32.279    0.000
placements           0.120      0.012  10.067    0.000
genreHipHop/Rap      0.199      0.048   4.188    0.000
genreRock           -0.256      0.060  -4.270    0.000
genreElectro/Dance   0.103      0.059   1.728    0.085
```

`log(spend)` = **0.33**: a 1% increase in spend goes with 0.33% more streams; +10% spend → **+3.3% streams** — the same *percentage* at every budget. Elasticities have no units and are **comparable**: followers (0.51) matter more than the budget.

This is a **marketing mix model** in miniature: sales on the marketing inputs, log-log form → an **elasticity per channel** → reallocate the budget toward the channels where a euro buys the largest percentage gain. Real MMMs add carryover (adstock) and saturation — and share this chapter's problem: budgets are not set at random, so spend is **endogenous**. Companies therefore calibrate MMMs with experiments.

Four ways to combine levels and logs



Model	Equation	Reading β_1
Level-level	$y = \beta_0 + \beta_1 x$	+1 unit $x \rightarrow \beta_1$ units of y
Log-level	$\log y = \beta_0 + \beta_1 x$	+1 unit $x \rightarrow \approx 100\beta_1$ % of y (exact: $100(e^{\beta_1} - 1)\%$)
Level-log	$y = \beta_0 + \beta_1 \log x$	+1% $x \rightarrow \approx \beta_1/100$ units of y
Log-log	$\log y = \beta_0 + \beta_1 \log x$	+1% $x \rightarrow \approx \beta_1$ % of y (**elasticity**)

Your turn (5 min)

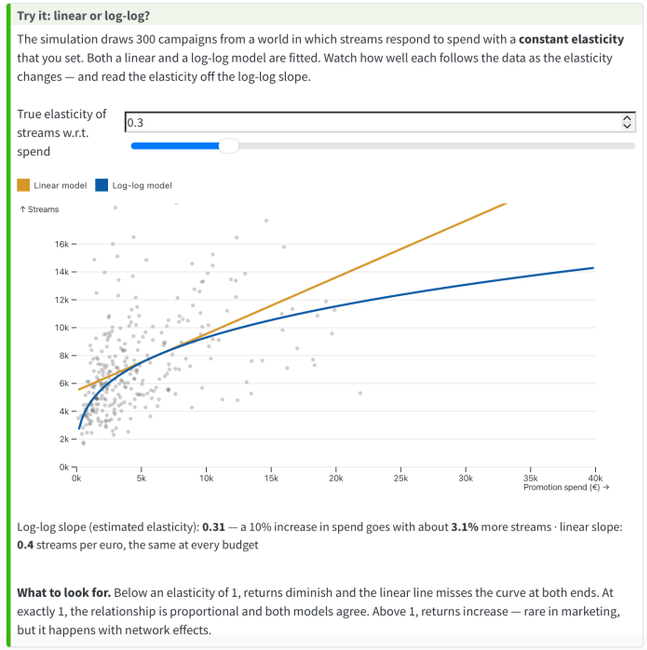
From the log-log model:

1. `genreHipHop/Rap` = 0.20. By what percentage do HipHop/Rap campaigns differ from Pop campaigns (the reference) — approximately and exactly?
2. A marketing manager proposes to double the promotion budget of every campaign. By roughly what percentage would streams increase, according to the model? Is this a causal claim?

Live demo: linear or log-log?

Course book, §5.7 — move the **true elasticity**.

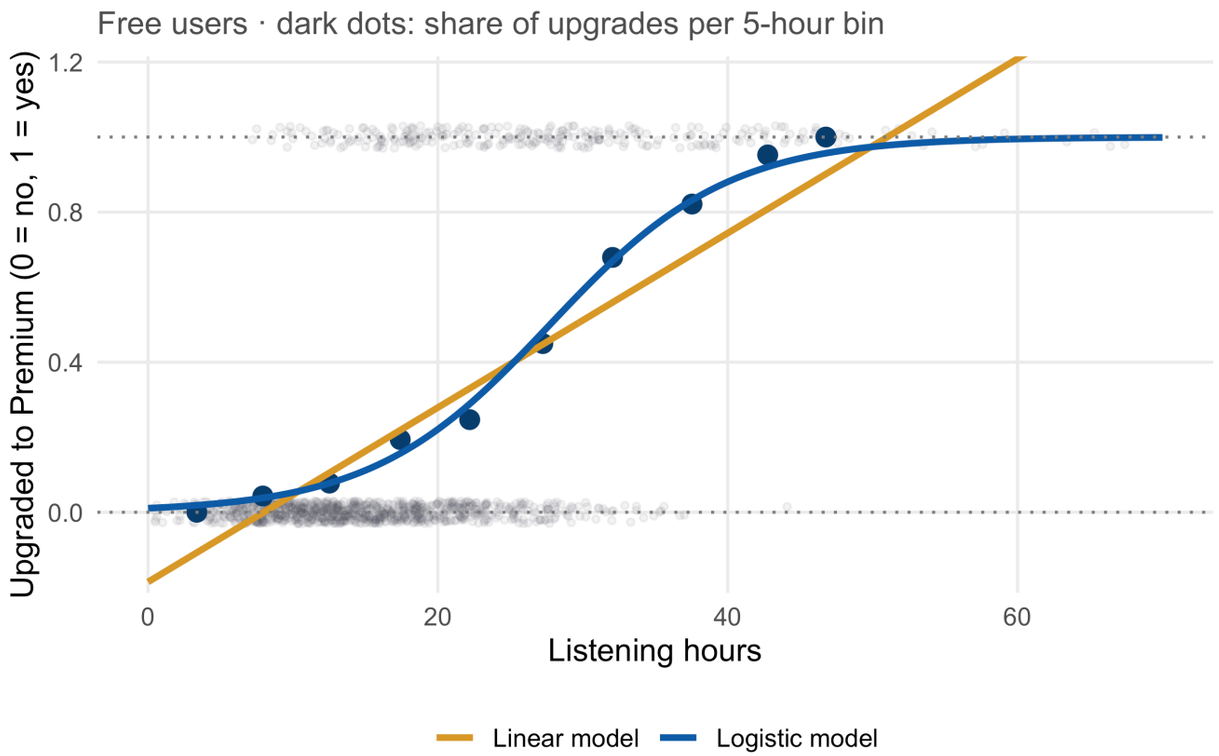
https://wu-rds.github.io/MA2026/05-observational-data.html#sec-obs-form



> Elasticity **< 1**: diminishing returns, the line misses at both ends. **= 1**: proportional. **> 1**: increasing returns (rare — network effects).

5.8 Yes/no outcomes: logistic regression

Why not a straight line?



- Outcome: **upgraded to Premium** within 3 months — free users, 22% did
- A line predicts probabilities **below 0** and **above 1** — and misses the shape
- The share of upgrades rises slowly, then steeply, then flattens near 100%: an **S-curve**

> We need a function that only takes values between 0 and 1.

The logistic function

$$P(y_i = 1) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_i)}}$$

- $\beta_0 + \beta_1 x_i$ is the familiar line — fed through a function that squeezes it into [0, 1]: line far **below** 0 → $P \approx 0$ · line = 0 → P = 0.5 · line far **above** 0 → $P \approx 1$
- Derived from the **log-odds**: odds = $P/(1 - P)$ ("1 to 4" = 0.25)

[Math Processing Error]

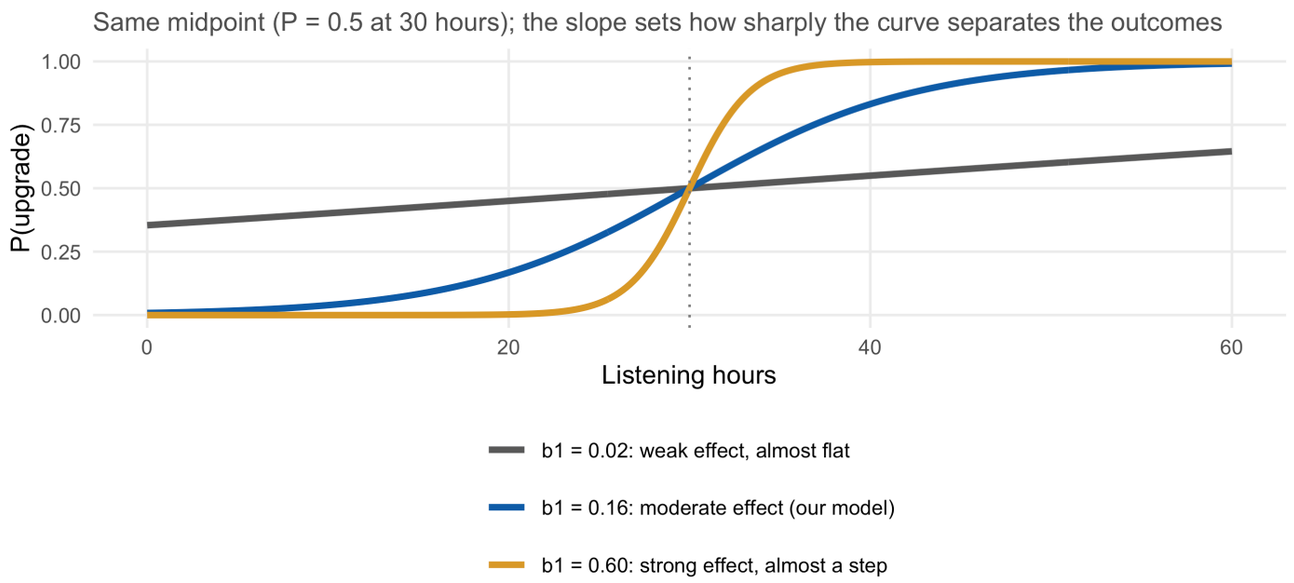
- Exponentiate → the odds: $\frac{P(y_i = 1)}{1 - P(y_i = 1)} = e^{\beta_0 + \beta_1 x_i}$

```r
free_users <- filter(smartmix_obs, premium == 0)
model_upgrade <- glm(upgraded ~ hours, family = binomial,
                     data = free_users)
round(coef(summary(model_upgrade)), 3)
```

```
            Estimate Std. Error z value Pr(>|z|)
(Intercept)   -4.517      0.244  -18.53        0
hours          0.163      0.011   15.05        0
```

> `glm()` = *generalized* linear model: the same predictors, a different link to the outcome. Estimated by **maximum likelihood**, not least squares (no line, no residuals to square): the coefficients under which the observed 0/1 data are most likely.

Weak vs. strong effects: the slope



- $\beta_1 \approx 0$: almost flat — the predictor hardly **discriminates** between upgraders and non-upgraders
- β_1 **large**: an almost vertical step — listening time nearly determines who upgrades
- β_0 shifts the curve; **midpoint** ($P = 0.5$) at $x = -\beta_0/\beta_1 \approx$ **28 hours**

> **Sign** = direction, **size** = how sharply the outcome responds — on the log-odds scale; “how much” in probabilities: next two slides.

Reading the coefficients: odds ratios

```
1 exp(coef(model_upgrade))
```

```
(Intercept)       hours
    0.01093     1.17689
```

```
1 exp(confint(model_upgrade))
```

```
              2.5 %  97.5 %
(Intercept) 0.006661 0.01734
hours       1.152949 1.20296
```

- Coefficient 0.163 = change in the **log-odds** per hour — no intuition
- $e^{0.163}$ = **1.18**: each hour multiplies the **odds** of an upgrade by 1.18 (+18%) — the same factor everywhere
- Read **sign** and **significance** directly from the output

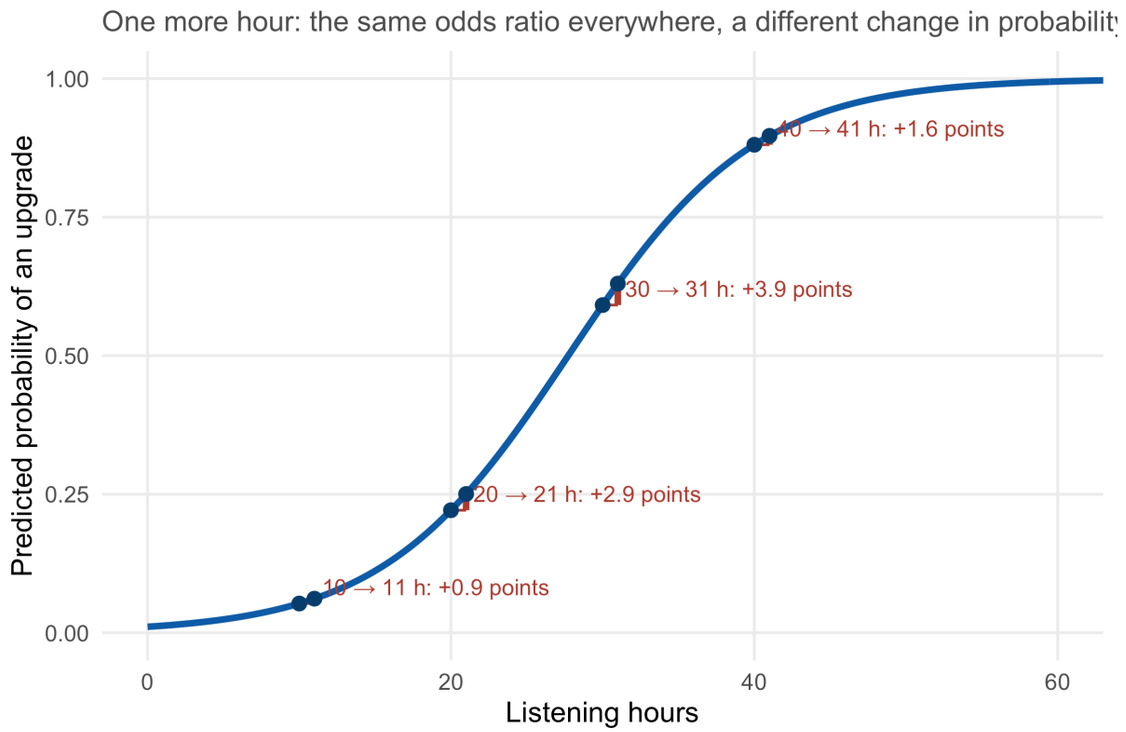
> **Odds ratio ≠ probability ratio.** “Odds +18%” is not “18% more likely” — unless the event is rare. For managers, translate into **probabilities**.

Reading the coefficients: probabilities at x and x + 1

```
predict(model_upgrade, type = "response",
        newdata = data.frame(hours = c(10, 11, 20, 21, 30, 31, 40, 41)))
```

```
      1       2       3       4       5       6       7       8
0.05275 0.06151 0.22111 0.25043 0.59135 0.63005 0.88062 0.89671
```

- 10 → 11 hours: **+0.9** points · 20 → 21: **+2.9** · 30 → 31: **+3.9** · 40 → 41: **+1.6**
- Same coefficient, same odds ratio — a **different change in probability**, largest around 50%



> Report effects as **predicted probabilities at chosen values** (10 vs. 30 hours: 5% vs. 59%) or as the **average change** across all users — “the effect of one more hour” has no single answer in percentage points.

Several predictors: churn

```
model_churn <- glm(churn ~ hours + premium + tenure + smart_mix, family = binomial, data = smartmix_obs)
round(coef(summary(model_churn)), 3)
```

```
            Estimate Std. Error z value Pr(>|z|)
(Intercept)    0.277      0.161   1.717    0.086
hours         -0.061      0.009  -7.130    0.000
premium       -0.533      0.153  -3.489    0.000
tenure        -0.020      0.004  -5.452    0.000
smart_mix     -0.256      0.162  -1.579    0.114
```

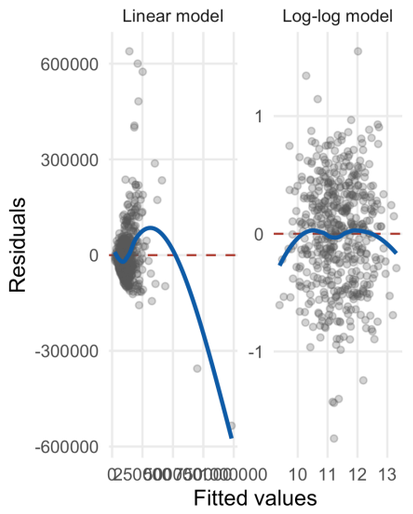
```
p_free    <- predict(model_churn, newdata = mutate(smartmix_obs, premium =
p_premium <- predict(model_churn, newdata = mutate(smartmix_obs, premium =
mean(p_premium - p_free)
```

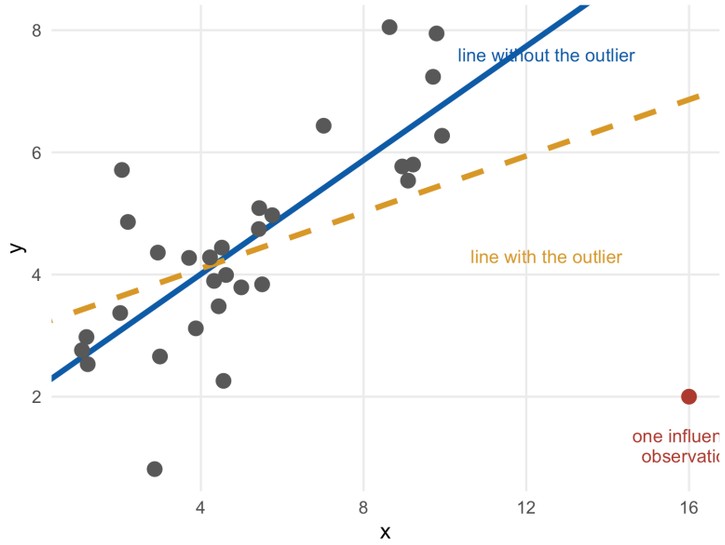
```
[1] -0.06477
```

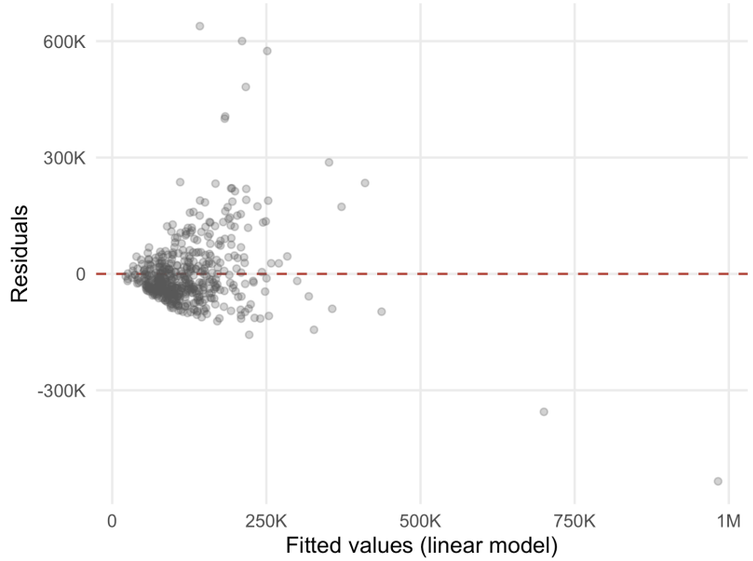
Same logic as multiple regression: each coefficient holds the others constant. Premium subscribers: about **9 percentage points** less likely to churn than otherwise identical free users. Causal? Exogeneity, again.

5.9 The diagnostics that change decisions

Three checks worth doing every time







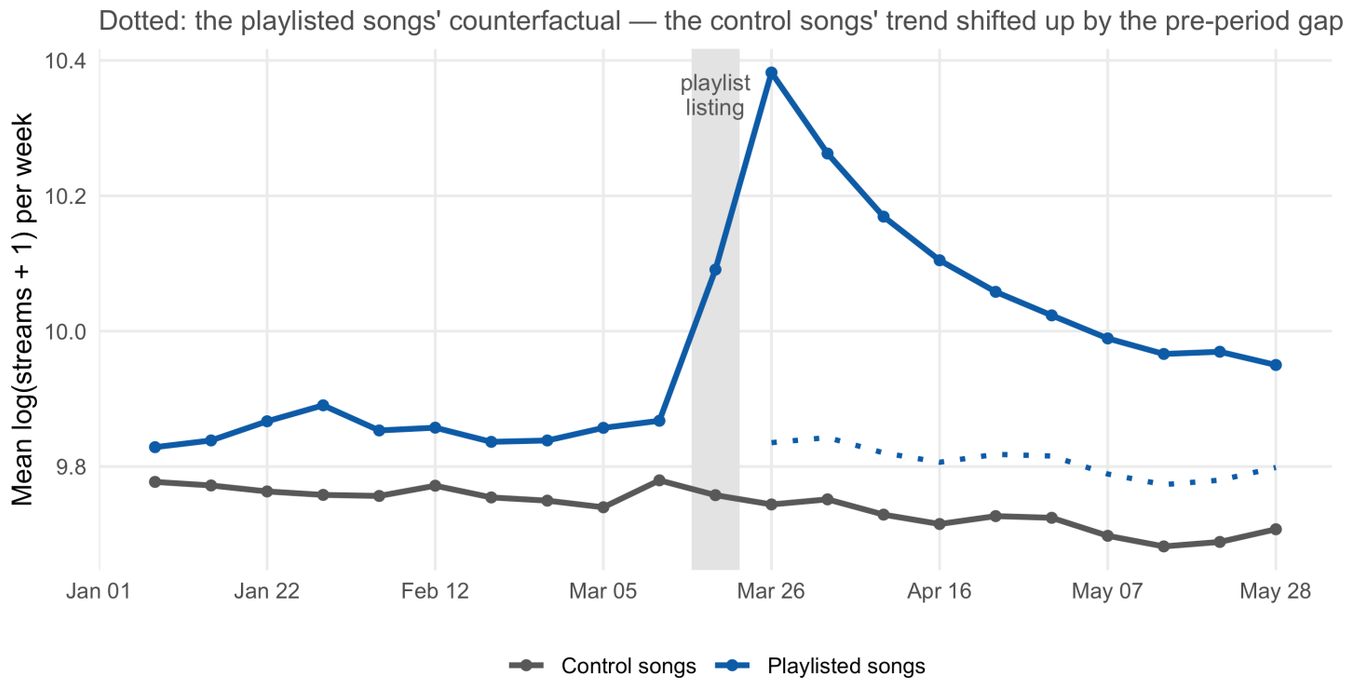
1. Is the line the right shape? Curved residuals → change the functional form (log) or add a predictor.

2. Does one observation drive the result? `plot(model, which = 5)`; refit without it — if the conclusion changes, say so.

3. Is the uncertainty right? A funnel → **robust standard errors**: `coeftest(model, vcov = vcovHC(model))`.

5.10 A quasi-experiment: difference-in-differences

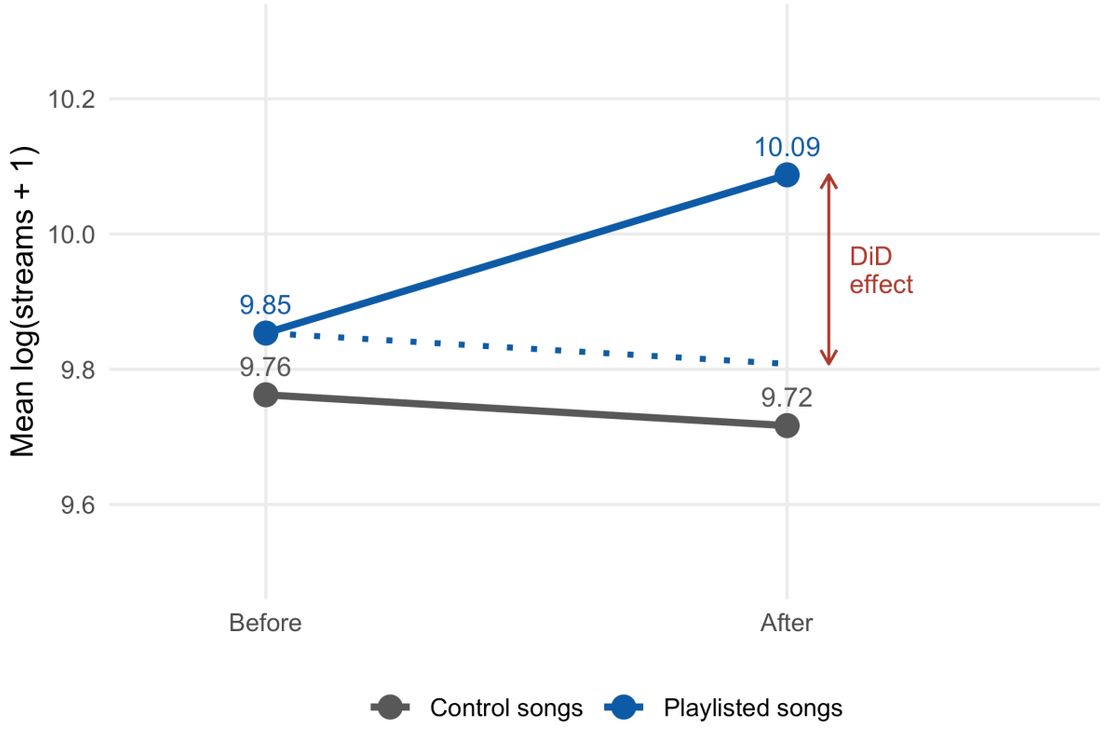
Playlist listings as a natural experiment



- **Real data:** 290 songs × 21 weeks
- March 19, 2018: **53 songs** added to a major playlist; 237 control songs
- Treated songs start **higher** → levels not comparable
- Before the listing: both groups move in **parallel**
- At the listing: the treated group **jumps**

Data: WU-RDS research project (github.com/WU-RDS/did).

Two differences, one estimate



$$\text{DiD} = (\bar{y}_{T,\text{after}} - \bar{y}_{T,\text{before}}) - (\bar{y}_{C,\text{after}} - \bar{y}_{C,\text{before}})$$

(10.09 − 9.85) − (9.72 − 9.76) = **0.28** log points

> The control group's change estimates what would have happened to the treated songs **without** the listing — the counterfactual (dotted).

The DiD regression

```r
did_data <- playlist |>
  filter(week != as.Date("2018-03-19")) |>          # drop the listing week
  mutate(log_streams = log(streams + 1))
model_did <- lm(log_streams ~ treated * post, data = did_data)
round(coef(summary(model_did)), 3)
```

```
            Estimate Std. Error t value Pr(>|t|)
(Intercept)     9.762      0.018 545.865    0.000
treated         0.091      0.042   2.182    0.029
post           -0.046      0.025  -1.803    0.071
treated:post    0.280      0.059   4.727    0.000
```

Coefficient	Meaning
`(Intercept)` 9.76	Control, before: baseline
`treated` 0.09	Treated, before: the **level difference**
`post` −0.05	Control, after − before: the **common time trend**
`treated:post` **0.28**	The extra change of the treated songs: the **DiD effect** → ≈ **+32% streams**

The assumption: parallel trends

```r
placebo_data <- did_data |>
  filter(post == 0) |>
  mutate(fake_post = ifelse(week >= as.Date("2018-02-12"), 1, 0))
round(coef(summary(lm(log_streams ~ treated * fake_post,
                      data = placebo_data)))["treated:fake_post", ], 3)
```

```
Estimate Std. Error    t value   Pr(>|t|)
   0.002      0.085      0.027      0.978
```

> Without the listing, the treated songs would have followed the **same trend** as the controls. **Untestable** — but the pre-period makes it plausible: a fake listing in the pre-period shows **no effect**.

Live demo: course book, §5.10

https://wu-rds.github.io/MA2026/05-observational-data.html#sec-obs-did

Your turn (5 min)

Add `major_label` as a control variable to the DiD regression:

```
summary(lm(log_streams ~ treated * post + major_label, data = did_data))
```

Does the effect change? Then: name one situation in which the parallel-trends assumption would **fail** for playlist listings.

Coding block: two cases (40 min)

Case 3: podcast promotion, revisited — `podcast_promo.csv` again.

1. Scatter plot with **both axes logged**: what has changed?
2. **Log-log model** on spend and host followers; interpret both **elasticities**: what does +10% spend do?
3. Compare $\mathbf{R^2}$ with the linear model; which model for the budget advice?

Laptops open — alone or with your neighbor. Section "Coding block" in `05-observational-data.R`.

Case 4: winning back users — `app_reinstall.csv`, 1,200 users, **randomized** discount 0–60%, reinstalled (1/0), prior sessions.

1. **Share of reinstalls per discount level**, plotted: describe the shape
2. **Logistic regression** of `reinstalled` on `discount` and `prior_use`: sign and **odds ratio** of the discount
3. Predicted probabilities for 10 prior sessions at **10, 11, 40, 41%**: points per extra point of discount — why do the two answers differ?

> At the end: one table presents case 3, another case 4 — Menti first: your spend elasticity (case 3) and your odds ratio (case 4).

What observational data can and cannot tell you

Two strategies, one honest report

- **Adjustment** handles confounders you can **measure** — only as good as your proxies; “enough” is untestable
- **Difference-in-differences** handles confounders **constant over time** — needs parallel trends, plausible but never proven
- **Prediction** needs no causal assumptions — a good predictor is valuable even if no coefficient is an effect

> Report effect + uncertainty + **assumptions**: “Adjusting for prior listening, tenure, Premium and device, adopters listen ≈ 4.5 hours more (95% CI 3.6–5.3), assuming no further differences affect listening — the experiment (3.2) suggests this is not fully met.”

Assignment 3: analyzing observational data

- Your **personal** data (based on your student ID): an opt-in feature, promotion campaigns, and a small panel with a listing
- Tasks: naive vs. adjusted · categorical predictor · elasticity · logistic regression · DiD
- **Deadline:** Group A Sun, Nov 22 · Group B Wed, Nov 25
- **Debrief** at the start of the following session

10 points: report 7 + debrief 3. State your **assumptions** — the best answers say what the data *cannot* show.

Learning check (Part II)

1. Log-log coefficient of ad spend = 0.15. Interpret. The level model says 0.8 sales per euro — why might the two disagree about doubling the budget?
2. Log-level model, dummy “promotion week” = 0.35. Percentage difference in sales — approximate and exact?
3. Residual plot shows a funnel. Consequence for coefficients, for *p*-values — and the simplest fix?
4. A DiD of a price promotion in treated stores: which assumption, and how can the pre-period help?

Next session

Module 6 · Mechanisms & heterogeneity — *Why, and for whom, does it work?*

- Interaction effects: when the effect differs between groups (moderation)
- Mediation: through which mechanism an effect works
- In experiments and observational data — and in your group project