

Marketing Analytics

Chapter 6 · Mechanisms & Heterogeneity

Nils Wlömert

Vienna University of Economics and Business (WU)

October 2026

Before we start

Paper quiz on Chapter 5, Part II (§5.7–5.10, 8 min)

- Individual, closed book
- Counts toward participation

Assignment 3 is running

- Deadline: Group A Sun, Nov 22 · Group B Wed, Nov 25
- Debrief at the start of next session

Today's session

- 6.1 The average effect hides the interesting part
- 6.2 Interactions and **marginal effects**
- 6.3 Continuous $\times$ categorical
- 6.4 Continuous $\times$ continuous (*short*)
- 6.5 Categorical $\times$ categorical: a **2 $\times$ 3 experiment**
- 6.6 Pitfalls: subgroup fishing & power
- 6.7 **Mediation**: why does it work?
- 6.8 Good and bad controls (*short*)
- 6.9 Your group project

Driving question: *Why — and for whom — does it work?*

One new ingredient: the **product of two variables**.

6.1 The average effect hides the interesting part

Smart Mix, second wave

- 1,200 users, randomized to Smart Mix or control
- This time with **covariates**: plan (free/Premium), prior listening, tenure, device
- And one more variable: **new artists discovered** during the four weeks
- Average effect: **+3.7 hours** (95% CI 3.1–4.3)

The product team asks two more questions:

1. **For whom** does it work — free or Premium users? light or heavy listeners?
2. **Why** does it work — does it make users discover new music?

Moderation (for whom): the effect *depends on* a third variable. **Mediation** (why): the effect *runs through* a third variable. Easy to confuse — a moderator changes *how large*, a mediator is *part of the path*.

6.2 Interactions and marginal effects


What an interaction term does

$$y = \beta_0 + \beta_1 x + \beta_2 z + \beta_3 (x \times z) + \varepsilon \quad \Rightarrow \quad \frac{\partial y}{\partial x} = \beta_1 + \beta_3 z$$

- The **marginal effect** of x depends on z
- β_1 = effect of x **when** $z = 0$ (a *conditional* effect, not a “main effect”)
- β_3 = how much the effect of x changes per unit of z
- Symmetric: effect of $z = \beta_2 + \beta_3 x$

Two rules: (1) if zero is not a meaningful value of z , β_1 is not meaningful either $\rightarrow$ center (§6.4). (2) **Always keep both main terms:** in R, $x * z = x + z + x:z$.

Parallel, diverging, crossing




With a group indicator as z : **no interaction** = parallel lines (same slope, different level) · **interaction** = different slopes · **crossover** = the effect of x reverses. β_3 is simply the **difference between the slopes**.

6.3 Continuous × categorical

Face 1: a treatment effect that differs between groups


```
1 model_plan <- lm(hours ~ smart_mix * premium + prior_hours + tenure, data = wave2)
2 round(coef(summary(model_plan))[, c(1, 2, 4)], 3)
```

	Estimate	Std. Error	Pr(> t)
(Intercept)	1.560	0.469	0.001
smart_mix	4.999	0.378	0.000
premium	2.685	0.431	0.000
prior_hours	0.752	0.017	0.000
tenure	0.028	0.009	0.001
smart_mix:premium	-3.111	0.590	0.000



smart_mix = effect for free users (premium = 0): 5.0 h ·
 smart_mix:premium = how much the effect **differs** for Premium:
 -3.1 h (the heterogeneity test) · effect for Premium = 5.0 - 3.1 =
 1.9 h · premium = plan difference in the control group.

Both effects with their uncertainty




- The CI of the Premium effect is **not** in the output — it combines two coefficients and their covariance (`emmeans` does it for you)
- Actionable: promote Smart Mix to **free users first**

The interaction coefficient tests whether the effects **differ**; the two marginal effects tell you **what they are**. Report both.

Face 2: a slope that differs between groups

```
1 model_genre <- lm(log(streams) ~ log(spend) * genre + log(followers) + plac
2                               data = promo)
3 round(coef(summary(model_genre))[, c(1, 4)], 3)
```

	Estimate	Pr(> t)
(Intercept)	-3.648	0.000
log(spend)	0.206	0.000
genreHipHop/Rap	-2.058	0.000
genreRock	0.356	0.587
genreElectro/Dance	0.021	0.973
log(followers)	0.507	0.000
placements	0.122	0.000
log(spend):genreHipHop/Rap	0.268	0.000
log(spend):genreRock	-0.053	0.504
log(spend):genreElectro/Dance	0.005	0.946



Same model, different question: $\log(\text{spend})$ = elasticity for **Pop** (reference); $\log(\text{spend}) : \text{genreHipHop/Rap} = +0.27 \rightarrow$ HipHop/Rap elasticity ≈ 0.47 . Rock and Electro: no evidence of a different slope. Genre “main effects” = level differences at 1 euro of spend.


Moderation by a continuous variable

```
1 model_prior <- lm(hours ~ smart_mix * prior_hours + premium + tenure, data = data)
2 round(coef(summary(model_prior))[, c(1, 4)], 3)
```

	Estimate	Pr(> t)
(Intercept)	1.446	0.007
smart_mix	5.391	0.000
prior_hours	0.797	0.000
premium	1.041	0.001
tenure	0.027	0.002
smart_mix:prior_hours	-0.096	0.005

Marginal effect of Smart Mix = $5.39 - 0.096 \times \text{prior hours}$:

- 10 hours before → 4.4 h · 30 hours → 2.5 h
- **smart_mix** alone = effect at **zero** prior hours — a user who hardly exists



With a continuous moderator, report the effect at **meaningful values** (10th / 50th / 90th percentile) or as a **curve with its confidence band** — never the coefficient at zero.

Your turn (5 min)

Fit the interaction of Smart Mix with **device**:

```
1 model_device <- lm(hours ~ smart_mix * device + prior_hours + premium + tenure, data = wave2)
2 coef(summary(model_device))
```

For which device is the effect of Smart Mix largest? Compute the effect for each device from the coefficients. Are the differences significant?

6.4 Continuous × continuous

Spend × followers: the problem with “main effects”

Uncentered:

```
1 m_cc <- lm(log(streams) ~ log(spend) * log(followers) +
2           placements + genre, data = promo)
3 round(coef(summary(m_cc))[2:3, c(1, 4)], 3)
```

	Estimate	Pr(> t)
log(spend)	-0.896	0.000
log(followers)	-0.422	0.003

A **negative** spend elasticity? It is the elasticity for an artist with **one follower** (log = 0).

Centered at the means:

```
1 promo_c <- promo |> mutate(ls = log(spend) - mean(log(spend)),
2                             lf = log(followers) - mean(log(followers)))
3 m_ccc <- lm(log(streams) ~ ls * lf + placements + genre, data = promo_c)
4 round(coef(summary(m_ccc))[c(2, 3, 8), c(1, 4)], 3)
```

	Estimate	Pr(> t)
ls	0.299	0
lf	0.508	0
ls:lf	0.113	0


Now **ls** = elasticity for an **average** artist; the interaction is unchanged.

Centering changes the main effects (now “at the mean of the other variable”), **not** the interaction or the fit. Communicate with a **plot of the marginal effect** — book §6.4 — not with the table.

6.5 Categorical $\times$ categorical: a 2×3 experiment

Are AI-made ads judged as creative as human-made ads?

- 360 respondents rate the creativity of an ad (0–100)
- Factor 1 — **source** of the design: Human · AI · Human + AI
- Factor 2 — **disclosure**: source revealed · not revealed
- $2 \times 3 = 6$ cells, 60 respondents each, randomized



Revealed: human designs win. Not revealed: AI and hybrid designs are rated **as creative** — slightly higher. **Crossing lines = the effect of the source depends on disclosure.**

The regression: six coefficients, six cell means

```
1 model_ads <- lm(creativity ~ source * disclosure, data = ads)
2 round(coef(summary(model_ads))[, c(1, 2, 4)], 2)
```

	Estimate	Std. Error	Pr(> t)
(Intercept)	61.80	1.85	0
sourceAI	-13.53	2.62	0
sourceHuman + AI	-18.25	2.62	0
disclosureNot revealed	-9.67	2.62	0
sourceAI:disclosureNot revealed	22.23	3.71	0
sourceHuman + AI:disclosureNot revealed	28.15	3.71	0

Coefficient

Meaning (reference: Human, Revealed)

(Intercept) 61.8

Human designs, source revealed

sourceAI -13.5 · sourceHuman + AI -18.3

AI / hybrid designs, *when revealed*, vs. human

disclosureNot revealed -9.7

Human designs: not revealed vs. revealed

sourceAI:... +22.2 · sourceHuman + AI:... +28.2

The **extra** gain of AI / hybrid designs when not revealed, beyond the two main effects

Step 1: is there an interaction at all?

```
1 anova(model_ads)
```

Analysis of Variance Table

Response: creativity

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
source	2	1055	527	2.56	0.079	.
disclosure	1	4572	4572	22.17	3.6e-06	***
source:disclosure	2	13218	6609	32.04	1.6e-13	***
Residuals	354	73021	206			

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The ANOVA table tests each factor and the **interaction as a whole** (two coefficients at once). Interaction highly significant → the effect of the source depends on disclosure → **do not interpret the main effects on their own**. (Unbalanced cells: `car::Anova(model_ads, type = 2).`)

Step 2: which cells differ? Estimated marginal means

```
1 library(emmeans)
2 emm <- emmeans(model_ads, ~ source | disclosure)
3 pairs(emm)
```

disclosure = Revealed:

contrast	estimate	SE	df	t.ratio	p.value
Human - AI	13.53	2.62	354	5.161	<.0001
Human - (Human + AI)	18.25	2.62	354	6.960	<.0001
AI - (Human + AI)	4.72	2.62	354	1.799	0.1715

disclosure = Not revealed:

contrast	estimate	SE	df	t.ratio	p.value
Human - AI	-8.70	2.62	354	-3.318	0.0029
Human - (Human + AI)	-9.90	2.62	354	-3.775	0.0005
AI - (Human + AI)	-1.20	2.62	354	-0.458	0.8910

P value adjustment: tukey method for comparing a family of 3 estimates


From “there is an interaction” to “**where exactly** are the differences”: revealed → human beats AI (14 points) and hybrid; not revealed → AI and hybrid beat human by 9–10 points. *p*-values **Tukey-adjusted** within each disclosure condition — the two-step logic of Chapter 4, with one extra dimension.

Your turn (5 min)

1. Rebuild the cell mean for **hybrid** designs with the source **not revealed** from the coefficients of `model_ads`. Check it against the table of cell means.
2. Which **single number** in the regression output answers: “does disclosure change the rating of hybrid designs *more* than it changes the rating of human designs?”

6.6 The pitfalls of looking for heterogeneity


Subgroup fishing



- Plan × device × tenure × age × country ... = dozens of subgroups
- Testing each = **multiple testing** (Chapter 4)
- 20 subgroups, no heterogeneity: **64%** chance of a “significant” one
- Live demo: subgroup-fishing widget, book \$6.6

“Smart Mix works particularly well for Android users aged 25–34” after twenty subgroups = **noise with a story**.

Interactions need large samples



- SE of a *difference between two effects* = $2 \times$ SE of the overall effect
- Same power $\rightarrow 4 \times$ **the sample** for an interaction as large as a main effect
- And interactions are usually **smaller**
- Underpowered: miss real heterogeneity, **exaggerate** what you find

Pre-register the moderators you expect and test those. Everything else is exploratory — a hypothesis for the *next* experiment. Not significant $\rightarrow$ “no evidence that the effect differs”, not “the effect is the same”.

6.7 Mediation: why does it work?

Total = direct + indirect

Three regressions

```
1 model_m <- lm(new_artists ~ smart_mix + prior_hours + premium, data = wave2)          # path a
2 model_y <- lm(hours ~ smart_mix + new_artists + prior_hours + premium + tenure, data = wave2) # paths b and c'
3 a <- coef(model_m)["smart_mix"]; b <- coef(model_y)["new_artists"]; c_direct <- coef(model_y)["smart_mix"]
4 c(a = unname(a), b = unname(b), direct = unname(c_direct), indirect = unname(a * b), total = unname(c_direct + a * b))
```

a	b	direct	indirect	total
3.6185	0.3909	2.3063	1.4146	3.7209

- **a** = 3.6 more new artists with Smart Mix
- **b** = 0.39 hours per new artist, holding Smart Mix constant
- **c'** = 2.31 hours direct — down from the total 3.72

Indirect effect $a \times b = 1.42 \text{ hours} = 38\%$ of the effect runs through discovery. Discovery is **an** important mechanism — not the only one.

Uncertainty of the indirect effect: bootstrap

```

1 set.seed(1)
2 boot_indirect <- replicate(2000, {
3   d_boot <- wave2[sample(nrow(wave2), replace = TRUE), ]
4   coef(lm(new_artists ~ smart_mix + prior_hours + premium, data = d_boot))["smart_mix"] *
5     coef(lm(hours ~ smart_mix + new_artists + prior_hours + premium + tenure, data = d_boot))["new_artists"]
6 })
7 quantile(boot_indirect, c(0.025, 0.975))

```

```

2.5% 97.5%
0.9861 1.8632

```

A product of two estimates is **not normally distributed** → no textbook standard error → **bootstrap** (Chapter 3): resample users, re-estimate, look at the distribution. Report **total, direct, indirect** with uncertainty, plus the **share mediated**. (mediation package: `mediate()` does this incl. logistic outcomes.)

The catch: nobody randomized the mediator

Live demo: book §6.7 — mediation widget, switch on the **unobserved confounder of M and Y**

<https://wu-rds.github.io/MA2026/06-mechanisms.html#sec-mech-mediation>

Screenshot placeholder: the mediation widget
`img/ch06/widget_mediation.png`


Randomization identifies the **total** effect. The **split** needs one more assumption: **no unobserved confounder of mediator and outcome** (curious users discover more *and* listen more). Untestable — and the same problem as Chapter 5, one step down the chain.

What helps: measure likely confounders of M; **manipulate the mediator** (Smart Mix with vs. without discovery).

6.8 Good and bad controls

Fork, pipe, collider

Fork (confounder)




adjust for Z

Pipe (mediator)



do NOT adjust for Z
(unless you want the direct effect)

Collider (common effect)



do NOT adjust for Z

Fork = confounder → adjust (Ch. 5). **Pipe** = mediator → adjusting gives the *direct* effect, not the total. **Collider** = common effect → adjusting (or selecting on it) **creates** a spurious association. Draw the diagram **before** the analysis.

6.9 Your group project

Moderation and mediation with 120 respondents

Add a moderator?

- Only if central to your topic and you can state the **expected direction** in advance
- Measure it **before** the manipulation
- $\text{outcome} \sim \text{condition} * \text{moderator}$
- 60 per condition → only **large** interactions detectable
- A 2×2 spreads you over cells of 30

Measure a mediator?

- One or two plausible mechanisms, measured **after** the manipulation (perceived quality, trust, attention)
- Cheap — and tells you **where the chain broke** if the manipulation did not work
- Report total / direct / indirect **with the caveats**

State what was **planned** and what was **exploratory**. “No evidence that the effect differs” $\neq$ “the effect is the same”.

Learning check

1. `smart_mix` = 5.0, `smart_mix:premium` = -3.1. Effect for free users? For Premium? What does the interaction's p -value test?
2. 2×2 : intercept 40, `framingYearly` -5, `trialYes` +8, `framingYearly:trialYes` +6. The four cell means? Where does the trial help more?
3. A loyalty program's total effect on spending, with “store visits” as a control — fork, pipe, or collider? Consequence?
4. 25 segments tested, “works particularly well for returning customers on tablets” ($p = .02$). Problem — and what next?

Next session

Module 8 · Survey design — *from research question to questionnaire*

- Measuring constructs: items, scales, reliability
- Embedding your experiment in a Qualtrics survey
- Questionnaire structure, pretesting, data collection

Bring your group's **research question and planned manipulation** — we will work on your questionnaires in class.